

Dust Particles and Gas+Dust Self-Gravity in the AthenaK Shearing Box

IMPLEMENTATION DOCUMENT

Phase 4a: the dust-module merge; Phase 4b (i): the shear-periodic deposit fold; (ii): the ideal-gas energy update; (iii): the particle timestep under shear; Phase 4c: dust self-gravity; Phase 4d: the vertical policy, restart, diagnostics

`athenak-multigrid` tree, branch `dust-multigrid`

commits `449b8f09` (merge), `a7c0def3`, `774c6519`, `78320aee` (second merge), `6c0cafef`, `489c656c`, `375643e6`, `9a96f9f3`, d

September 3–5, 2026

Abstract

This document describes *how* Lagrangian dust particles are being brought into the `athenak-multigrid` tree — the shearing-box code with multigrid self-gravity on refined meshes documented in `multigrid_shear_implementation.pdf` — on the way to gas+dust self-gravity in a stratified, cooling, gravito-turbulent box (the configuration of Baehr, Zhu & Yang 2022, with the radial pressure gradient they omitted). It is the companion to `dust_selfgravity_validation.pdf`, which records what the resulting code does on test problems. Both documents are meant to grow phase by phase.

Phase 4a, recorded here, merges the sibling fork `athenak-dust` (particle-mesh dust with IMEX drag, validated in its own repository) into the multigrid tree. The merge itself is mechanical — eight keep-both conflicts and one renumbered output slot — but it exposed a real defect: the dust module’s additive ghost-deposit exchange was written against the per-neighbor MPI messaging of July’s upstream, while this tree carries the rank-packed messaging of upstream #775. Any multi-rank dust run deadlocked at cycle 0. The exchange now uses the aggregated-message path with the same idiom as the cell-centered exchange, and ranks 1, 2 and 4 agree to every printed digit. A second follow-up aligns the dust task graph with the tree’s Prolongate-before-boundary-conditions order.

A **second merge** the same day caught the tree up with the dust fork’s own latest work (seven commits, to `1f3e1ed9`): the fork had independently made the identical rank-packed rewrite of the deposit exchange, so that conflict resolves to their version; and it now offers two *predictor–corrector* coupling modes (`coupling = pc2` and the adaptive `hybrid`) that run under plain `rk2`, second order in the resolved drag regime with a stiff backward-Euler fallback. Dust can therefore run under the same integrator as every gravity validation before it; the `imex2+` path remains for stiff cases.

Phase 4b (i) closes the first of the gaps the survey named: ghost deposits of the MeshBlocks on the shear-periodic x_1 faces are now folded across the faces with a conservative azimuthal remap — gathered into global y – z planes, shifted by the shear offset with the Phase-1 remap-flux kernels, and *added* into the active edge strips of the blocks across the face — while the unsheared plain-periodic x_1 contributions of those blocks are suppressed. The fold is the adjoint of the copy-exchange ghost fill, so its shift signs are the reverse of the multigrid fill’s. A unit-test problem generator certifies it exactly at integer shifts and at second/third order otherwise, and a structured 3D run shows the old unsheared exchange was a percent-level error. **Phase 4b (ii)** lifts the isothermal restriction on back-reaction: every drag kick on the gas momentum now goes through one helper that, for an ideal gas, books the matching kinetic-energy change so

the internal energy is untouched, and a flagged option deposits the frictional dissipation of each particle kick as heat. The residual energy errors are the convexity gap of the low-storage RK combination, first order in the timestep, and are quantified. **Phase 4b (iii)** removes the last of the three Phase-4b gaps: the particle transport timestep no longer resolves the background shear $-q\Omega_0 x$ cell by cell (a $1/L_x$ penalty the orbitally advected gas does not pay) but limits the peculiar velocity per cell and the *total* transport velocity per MeshBlock — the latter being what the single-hop particle migration and the one-layer deposit halo actually require — with the bound made exact for the low-storage tableaus through their largest stage fraction. In an $L_x = 8$ box the step grows $5.3\times$ to the gas limit. **Phase 4c** adds the dust to the Poisson equation and the potential to the particles. The particle-mesh dust density is deposited before each stage’s solve, exchanged additively (with the shear fold) and then copied into its ghost layer, and registered with the gravity module as an *extra density* that both solvers add to the gas density at their single source read; after the solve the force $-\nabla\phi$ is formed on the mesh, exchanged like the module’s provisional gas velocity and gathered onto the particles with the deposit kernel. Because scatter and gather share one kernel and the gradient is the adjoint of the symmetric Laplacian, the coupling conserves the total gas+dust momentum to round-off by construction. Two limitations surfaced on the way and are recorded: the multigrid needs cubic cells and more than one MeshBlock across each periodic direction, and 3D runs on cubic cells grow a grid-scale mode once $c_s\Delta t/\Delta z \gtrsim 0.35$. **Phase 4d** closes the dust track’s remaining gaps. The instability just named is identified — by its eigenmode and by a seeded gas-only run — as the shearing-box hydro scheme’s own 3D checkerboard mode at Courant numbers above ~ 0.35 , not a property of the dust coupling, which merely seeded it; the rule $\text{cf1} \leq 0.3$ in 3D is a hydro rule the gravito-turbulence runs already follow. Particles crossing a physical vertical face are removed with their mass accounted (the Athena-C `zbc_out = 1` policy); particle arrays are written to and read from restart files; the module’s own particle-mesh density is an output variable; and the SC14 history gains the dust’s mass, momenta, kinetic energy, Reynolds stress and escaped mass.

The document also records, for the phases that follow, the integration map derived from the code survey: where the gravity solve is called relative to the dust task graph, the three places the solver reads density, the ghost-width contract of the potential, and the design chosen for the particle force (grid gradient, exchanged and shear-remapped like the module’s provisional gas velocity, gathered with the deposit kernel), together with the known gaps that Phases 4b–4d must close.

Contents

1	Scope, sources, and how to read this	5
1.1	What this document is	5
1.2	The program	5
1.3	Commit provenance	5
1.4	Files touched	7
2	The target problem and the design it implies	7
2.1	What is being built	7
2.2	The integration map	7
3	The dust module as inherited	8
3.1	Data layout	8
3.2	The particle-mesh kernel	9
3.3	The stage algorithm	9
3.4	The task graph	9

3.5	Shearing-box conventions	10
3.6	The particle timestep	10
3.7	Guards and parameters	10
4	Phase 4a: the merge	10
4.1	Route	10
4.2	Conflicts	11
4.3	Follow-up 1: task order	11
4.4	Follow-up 2: the deposit exchange on the rank-packed path	11
4.5	The second merge: catching up with the dust fork	12
4.6	Build notes	13
5	Phase 4b (i): the shear-periodic deposit fold	13
5.1	The problem	13
5.2	The sign	13
5.3	The fold	14
5.4	Suppressing the unsheared contribution	14
5.5	Hooks and parameters	15
5.6	The test generator	15
5.7	A structured, rank-independent initial condition	15
6	Phase 4b (ii): the ideal-gas energy update	15
6.1	The problem	15
6.2	The rule	16
6.3	What the bookkeeping cannot do, and the lesson from the first test	16
7	Phase 4b (iii): the particle timestep under shear	16
7.1	The problem	16
7.2	What the limit is actually for	17
7.3	The rule	17
7.4	Parameters and diagnostics	18
7.5	What the larger step exposes	18
7.6	What does not change	18
8	Phase 4c: dust self-gravity	18
8.1	Design, as built	18
8.2	Momentum conservation, exact by construction	19
8.3	Synchronous entry points	20
8.4	Parameters, guards, boundaries	20
8.5	What was found on the way	20
9	Phase 4d: the instability identified, the vertical policy, restart, diagnostics	21
9.1	The instability of §8 is the hydro scheme's	21
9.2	Vertical policy: outflow removal	21
9.3	Particle restart	21
9.4	Diagnostics	22
9.5	The science setup: Baehr, Zhu & Yang (2022) in two stages	22
10	What the next phases must add	22

1 Scope, sources, and how to read this

1.1 What this document is

This is an implementation reference for the dust track. For each phase it answers: what the code relies on, where that is established and enforced, which data structure carries the new state, which function does the work, and from which call site it is reached. Results live in the companion documents:

<code>validation/dust_selfgravity_validation.pdf</code>	results, tests, figures (this track)
<code>validation/multigrid_shear_implementation.pdf</code>	the gravity/refinement machinery being extended
<code>validation/multigrid_selfgravity_validation.pdf</code>	its validation record (Phases 0–3)
<code>../athenak_dust_tests/</code>	the dust module’s own validation notebooks and inputs

1.2 The program

The target science is the interaction of gravitational instability of the gas with dust concentration and collapse in a shearing box: Baehr, Zhu & Yang (2022, ApJ 933, 100) in 3D stratified, cooling, gravito-turbulent boxes with superparticles at $St = 0.1$ –10, plus the radial pressure gradient (and hence the streaming instability) that their setup left out, and eventually adaptive zoom on the dust clumps. The phases:

Phase	Status	Content
4a	done	merge athenak-dust into the multigrid tree; establish that both the dust module and the gravity machinery are unchanged by the merge
4b	(i), (ii) done	(i) shear-periodic <i>additive</i> deposit remap (§5); (ii) ideal-gas energy update under drag back-reaction (§6); (iii) particle timestep on the shear-subtracted velocity; adopt the PC2 coupling (§4.5) as the production path
4c	planned	dust self-gravity: mass deposit into the Poisson source, three density reads, force gather to particles; dust-only and gas+dust tests
4d	planned	the science configuration: SC14 box + pressure gradient + dust; particle restart; clump diagnostics
5	deferred	particles on refined meshes (deposit exchange across levels, particle redistribution on remesh)

1.3 Commit provenance

The merged fork is github.com/athenak-dust/athenak, branch `dust`. At the first merge it stood at `d6de1296`, thirteen commits on top of upstream AthenaK as of July 8, 2026 (commit `5f199310`, the merge base with this tree). The one that matters is `e785735e` “Add IMEX dust-particle coupling to hydrodynamics” (36 files, +3864 lines); the rest are the matrix-free coupled drag solvers, GPU migration fixes, a snapshot-seeded clump benchmark, and diagnostics. The local clone turned out to be two weeks stale; the second merge (§4.5) brought in the seven commits through `1f3e1ed9` (September 2). The remote is registered locally as `dust` (`git remote add dust ../athenak-dust`),

Commit	Phase	Content
449b8f09	4a	Merge of <code>dust/dust</code> (d6de1296) into <code>multigrid</code> (b4d4c547); eight conflicts resolved keep-both; output slot <code>grav_phi</code> 153 → 154
a7c0def3	4a	Dust task graph: <code>Prolongate</code> before <code>ApplyPhysicalBCs</code>
774c6519	4a	<code>bvals_dep.cpp</code> : additive deposit exchange on the rank-packed MPI vars path (multi-rank deadlock fix)
78320aee	4a	Second merge, <code>dust/dust</code> at 1f3e1ed9: the fork’s own rank-packed deposit exchange (taken), PC2/hybrid coupling, PM optimizations; three conflicts (§4.5)
6c0cafef	4b (i)	<code>MeshBoundaryValuesDep::FoldShearDeposit</code> , the x_1 -buffer skip, the dust hooks, the <code>dust_deposit_shear</code> test generator, the 3D shear-periodic NSH input, the NSH mass modulation
489c656c	4b (ii)	<code>GasKick</code> , the four kick sites, the heat channel of the back-reaction deposit, <code><dust>/drag_heating</code> , the isothermal guard removed, ideal-gas initialization of <code>dust_nsh</code>
375643e6	4b (iii)	<code>DustGasDrag::ComputeNewTimeStep</code> : cell limit on the peculiar velocity, <code>MeshBlock</code> guard on the total velocity; <code><dust>/dt_transport</code> , <code>dt_block_safety</code> ; <code>DUST_DT_SUMMARY</code>
9a96f9f3	4c	<code>dust/dust_gravity.cpp</code> ; <code>Gravity::RegisterExtraDensity</code> and the source read of both solvers; the dormant-stage solve skip; the force gather in <code>ExplicitPush</code> ; <code>Particles::dtnew</code> initialized; the twin and orbit generators, three inputs, NSH momentum history columns
d39858ac	4d	escape marking in <code>SetNewPrtclGID</code> , <code>Particles::RemoveDead</code> , the <code>RemoveEscaped</code> tasks; particle restart (writer step 3/5, reader); <code>dust_dpm</code> output; SC14 history dust columns; the seeded gas-only diagnostic (<code><problem>/gas_vpert</code>); <code>dust_grav_orbit vz0</code>
a9279dc8	science	the Baehr et al. two-stage setup: <code>bcool_cs2_floor</code> , <code>restart_insert</code> , <code>GravitoTurbInsertDust</code> , the two inputs, the Vista scripts

Table 1: Implementation commits covered by this document.

so later fixes in that fork come in with a plain `git merge dust/dust` — which is exactly what the second merge was.

1.4 Files touched

File	Role
src/dust/dust.{cpp,hpp}	DustGasDrag module: parameters, arrays, boundary objects, guards, the shared PM kernel <code>dust::PMWeights</code>
src/dust/dust_tasks.cpp	the combined hydro+particle task graph (<code>AssembleDustGasDragTasks</code>) and the boundary tasks
src/dust/dust_pm.cpp	deposit, per-cell implicit solve, gather/kick, back-reaction
src/dust/dust_push.cpp	explicit particle push (Coriolis, tide, drift, drag history)
src/dust/dust_newdt.cpp	particle transport timestep, stopping-time bookkeeping, the γ -switch
src/dust/dust_solver.cpp	optional matrix-free coupled drag solvers
src/bvals/bvals_dep.cpp	<code>MeshBoundaryValuesDep</code> : additive ghost-deposit exchange (rewired in 4a)
src/bvals/bvals_part.cpp, bvals.{cpp,hpp}	shear-periodic particle routing; deposit-exchange class declaration
src/driver/driver.{cpp,hpp}	<code>SetImEx2PlusCoefficients</code> (the γ -switch entry point)
src/mesh/meshblock_pack.{cpp,hpp}	<code>pdust</code> registration; hydro/particle task lists ceded to the dust module when a <code><dust></code> block exists
src/mesh/mesh.{cpp,hpp}, build_tree.cpp	<code><time>/dtmax</code> hard cap
src/athena.hpp	widened <code>ParticlesIndex</code> (17 real, 3 integer fields)
src/outputs/*	<code>dust_d</code> derived output; particle velocities in <code>pvtk</code> ; slot renumbering
src/pgen/tests/dust_*.cpp, streaming_linear.cpp	the four built-in dust problem generators
inputs/dust/*.athinput	their reference inputs

2 The target problem and the design it implies

2.1 What is being built

The reference implementation for particle self-gravity in a shearing box is the Athena-C problem generator `parsg_strat3d.c` in `athena_parsg/src/prob/` (Rixin Li's fork). Its method, with `PLUS_GAS_SELF_GRAVITY`: deposit the dust mass onto the grid with the TSC kernel, *add the gas density into the same buffer*, solve *one* Poisson equation, difference Φ on the grid, and interpolate the resulting acceleration to the particles with the *same* TSC kernel used for the deposit. The gas feels the total potential as a body force. Baehr et al. do the same in PENCIL (their eq. 15, $\rho = \rho_g + \rho_d$ in a periodic FFT). Two details of the Athena-C version are not to be copied: its deposit omits the cell-volume division and hides it in the particle mass, and its drag feedback must subtract the gravitational kick because gravity is folded into the same implicit particle update as drag.

2.2 The integration map

The survey of this tree fixed where each piece attaches. These are facts about the code as merged; the items marked *planned* are the Phase 4b/4c designs.

Where gravity is solved. Gravity has no task-list entry. The driver calls `Gravity::Solve(stage)` once per explicit stage, *between* the `before_stagen` and `stagen` task lists (`driver/driver.cpp`, lines 563–569). Anything the solver must see — the dust mass deposit and its ghost exchange — has to complete inside `before_stagen` (*planned*: a deposit task there, on the stage-start particle positions).

The three density reads. The solver reads gas density in three places, and all three must learn about dust (*planned*, Phase 4c):

- `gravity/mg_gravity.cpp:353`, `LoadSource(u0, IDN, ...)` into the multigrid source: loads active zones *plus one ghost layer*, so the dust deposit must already be additively ghost-exchanged one layer deep;
- `multigrid/multigrid_slab.cpp:263`, the slab-plane gather that builds the open- z Dirichlet planes (active zones, reduced across ranks) — omitting dust here would make the planes inconsistent with the interior source;
- `gravity/fft_gravity.cpp:218`, the FFT oracle’s density gather.

The potential’s ghost contract. `phi` is allocated with the hydro ghost width but only *one* ghost layer is ever filled (`mg_nghost = 1`; larger values are fatal with shear). The gas source term (`SourceTerms::SelfGravity`, `srcterms/srcterms.cpp:298`) needs exactly that: a centered $-\rho\nabla\Phi$ for momentum and the flux-based energy term of Mullen, Hanawa & Gammie (2020). A TSC gather at a particle in the last active cell reaches one ghost cell of the *acceleration*, i.e. two of the potential. *Planned*: compute $\mathbf{g} = -\nabla\Phi$ on active cells, then copy-exchange and shear-remap the three-component field exactly as the module already does for its provisional gas velocity u^* (`MeshBoundaryValuesCC` plus a `ShearingBoxCC`), and gather with `dust::PMWeights`. Using the deposit kernel for the force gather — not the derivative of the kernel — is the Hockney–Eastwood choice that suppresses self-force, and it is what the Athena-C reference does.

Where the particle force goes. `DustGasDrag::ExplicitPush` (`dust/dust_push.cpp:65`) already assembles the Coriolis and vertical-tide accelerations per particle; the gathered $\mathbf{g}$ joins them there. The back-reaction deposits only the drag kick $\Delta v_j = c_j(\tilde{u}_j^* - v_j)$, so gravity is never double-counted — there is no analogue of the Athena-C `feedback_corrector` subtraction to write.

The gas force. Nothing to add: the existing source term applies whatever potential the solver produced.

Stage structure. Under `coupling = imex`, `imex2+` has three explicit stages of which the third is dormant for explicit physics ($\beta_3 = 0$); the gravity solve should be skipped there or a third of the solver cost is wasted (*planned*). Under the PC2/hybrid couplings of §4.5 the integrator is `rk2` and the gravity solve runs exactly as it does for gas alone.

3 The dust module as inherited

This section is the working summary of `athenak-dust`, with the facts the later phases build on. The design is Yang & Johansen (2016) particle-mesh coupling (one shared kernel for scatter and gather, exact momentum-conserving back-reaction) with the Benítez-Llambay, Krapp & Pessah (2019) closed-form per-cell implicit drag inside the Krapp et al. (2024) IMEX(4,3,2) integrator (`imex2+`), generalized to per-particle stopping times.

3.1 Data layout

Particles are two flat device arrays indexed (`field`, `particle`). The field enumeration (`athena.hpp:76`) is


```
enum ParticlesIndex {PGID=0, PTAG=1, PSP=2,
                    IPX=0, IPVX=1, IPY=2, IPVY=3, IPZ=4, IPVZ=5,
                    IPX1=6, IPVX1=7, IPY1=8, IPVY1=9, IPZ1=10, IPVZ1=11,
                    IPRX=12, IPRY=13, IPRZ=14, IPTS=15, IPM=16};
```

so a dust particle carries position, velocity (all three components even in 2D r - z), the low-storage stage-1 registers, the recorded drag rate R_j , its stopping time and its **mass**. Integer data: owning MeshBlock, global tag, species. `nrdata = 17, nidata = 3`.

3.2 The particle-mesh kernel

`dust::PMWeights` (`dust/dust.hpp:78`) returns the three weights of a NGP / CIC / TSC stencil in one direction; every deposit, gather, matrix-free operator and kick uses it, which is what makes scatter and gather weights identical by construction. `<dust>/deposit` selects the kernel (default `tsc`); the stencil is always three wide (NGP zeroes the wings). Particle-mesh assignment needs `nghost >= 2`.

3.3 The stage algorithm

Per explicit stage, on the conserved gas state (primitives are stale mid-stage), with $a \Delta t = a_{\text{impl}} \Delta t$ the implicit weight of the tableau:

1. **DepositDrag**: scatter $Q = \sum_j \mu_j W_{jk}$ and $\mathbf{P} = \sum_j \mu_j W_{jk} \mathbf{v}_j$ with $c_j = a \Delta t / (t_{s,j} + a \Delta t)$, $\mu_j = m_j c_j / V$ (a *drag-weighted* density, not a mass density); additive halo exchange.
2. **GasImplicitSolve**: per cell, $\mathbf{u}^* = (\rho \mathbf{u} + \mathbf{P}) / (\rho + Q)$, the closed-form momentum-conserving backward-Euler pair. With `back_reaction = false` the deposits are excluded so test particles are kicked toward the unmodified gas.
3. Copy exchange of $\mathbf{u}^*$ ghosts (plus shear remap in a 3D shearing box).
4. **GatherKickPMBR**: gather $\tilde{\mathbf{u}}_j^*$, kick $\Delta \mathbf{v}_j = c_j (\tilde{\mathbf{u}}_j^* - \mathbf{v}_j)$, record $R_j = \Delta \mathbf{v}_j / (a \Delta t)$, and deposit $-m_j W_{jk} \Delta \mathbf{v}_j / V$ with the *same* weights and the same $\Delta \mathbf{v}$ — exact momentum conservation to round-off.
5. Additive halo exchange of that deposit; **ApplyPMBR** adds it to the gas momentum and rescales it in place into the gas drag rate R_g .
6. Next stage, the history terms $\tilde{a}_{22} \Delta t R$ enter on both sides (**AddDragHistoryGas**; the particle side inside **ExplicitPush**).

There is *no* drag timestep constraint for any t_s or dust-to-gas ratio; only particle transport limits Δt (§3.6).

3.4 The task graph

When a `<dust>` block exists, `MeshBlockPack::AddPhysics` skips both `Hydro::AssembleHydroTasks` and `Particles::AssembleTasks` and hands the whole graph to `DustGasDrag::AssembleDustGasDragTasks` (`mesh/meshblock_pack.cpp:265-271`, `dust/dust_tasks.cpp:37-108`), the same pattern as `IonNeutral`. The consequence for later phases:

Ownership. The dust module owns the entire per-stage task graph. Any task the gravity work needs (deposit, exchange, force gather) is inserted into `AssembleDustGasDragTasks`, and the hydro chain inside it must track changes made to `Hydro::AssembleHydroTasks` in this tree (the 4a follow-up `a7c0def3` is the first instance).

The **stagen** sequence: **FirstTwoImpRK** (replaces **CopyCons**) → gas fluxes / update / source terms → **AddDragHistoryGas**; in parallel **ExplicitPush** → particle migration (**NewGID**, counts, sends, receives — *inside every stage*, not once per cycle) → the deposit chain of §3.3 → the gas tail (orbital advection, restrict, exchange, shear exchange, **Prolongate**, **ApplyPhysicalBCs**, **ConToPrim**, timesteps). Every dust task returns early on the dormant third stage.

3.5 Shearing-box conventions

Particle velocities are *shear-subtracted*: the explicit push applies the FARGO-form Coriolis terms $f_x = 2\Omega_0 v_y$, $f_y = -(2 - q)\Omega_0 v_x$ and, with `<shearing_box>/stratified`, $f_z = -\Omega_0^2 z$ (`dust/dust_push.cpp:98-100`); positions drift with the *full* azimuthal velocity, $\dot{y} = v_y - q\Omega_0 x$ (:113-114). This matches the gas with `orbital_advection = true` (the tree’s default and the SC14 setting). Crossing a radial face is therefore a *position-only* transform: wrap x , shift y by $\mp y_{\text{sh}}$, $y_{\text{sh}} = q\Omega_0 L_x t \bmod L_y$ (`bvals/bvals_part.cpp:76,134`), with the RK position registers shifted by the same amounts so the low-storage combination stays in one periodic image. The destination **MeshBlock** is found from a precomputed $(y, z) \rightarrow (\text{gid}, \text{rank})$ map of the face blocks, not from the x_1 neighbor.

The radial pressure gradient is a gas-only constant acceleration through the stock `<hydro_srcterms>/const_accel` machinery ($2\Omega_0 \eta v_K$ along x); particles never feel it.

3.6 The particle timestep

`DustGasDrag::NewTimeStep` (`dust/dust_newdt.cpp:95`) takes $\min_p \min_i \Delta x_i / |v_{\text{transport},i}|$ times `<dust>/dt_cfl` (default 0.5), where in a 3D shearing box $v_{\text{transport},y} = v_y - q\Omega_0 x$ (:120). In a wide box this is the same penalty the gas pays without orbital advection ($\sim 5\times$ at SC14 size). Phase 4b (iii) relaxes it (§7).

3.7 Guards and parameters

Fatal conditions in `DustGasDrag::DustGasDrag` (`dust/dust.cpp`): MHD present (:45); 1D mesh (:50); an integrator other than `imex2+` (:63); `multilevel` (:69); boundaries neither strictly periodic nor 3D with shear-periodic x_1 (:78); `nghost < 2` (:93); back-reaction with a non-isothermal EOS (:170); a coupled solver other than `local` with shear (:227). A *warning* at :84 states that ghost deposits are exchanged with the plain-periodic pattern under shear, exact only for azimuthally uniform states. The same `multilevel` guard is repeated in the `MeshBoundaryValuesDep` constructor and in the shear-periodic particle constructor. The parameter interface is reproduced in §11.

4 Phase 4a: the merge

4.1 Route

The two forks share upstream history through 5f199310 (July 8, 2026); the dust fork adds 13 commits, this tree about 140 including an upstream merge on August 10 that brought in #775 (rank-packed boundary messaging) and #612 (`imex2+`). A merge was preferred over cherry-picks: it keeps the dust fork’s history and authorship, resolves conflicts once, and makes later fixes from that fork a one-line pull. The whole sequence was first rehearsed in a scratch clone (conflicts, follow-ups, both builds, the full acceptance battery) and then repeated in the working tree with the rehearsed resolutions copied in.

4.2 Conflicts

Eighteen files were touched by both forks; eight conflicted, all of the same kind — both sides appended adjacent code at the same anchor — and all were resolved by keeping both sides:

File	this tree had added	the dust fork had added
<code>driver/driver.cpp</code>	the super-time-stepping controller functions	<code>SetImEx2PlusCoefficients</code> (the constructor’s inline <code>imex2+</code> tableau moved into a function the γ -switch can call)
<code>mesh/mesh.cpp</code> , <code>mesh.hpp</code> <code>mesh/meshblock_pack.cpp</code> , <code>.hpp</code>	STS timestep cap and members physics (9) gravity, <code>pgrav</code>	<code>dtmax</code> hard cap (default ∞ , inert) physics (10) dust, <code>pdust</code>
<code>outputs/outputs.hpp</code> , <code>basetype_output.cpp</code>	<code>grav_phi</code> at slot 153 and its guard	<code>dust_d</code> at slot 153 and its guard
<code>parameter_input.cpp</code>	gravity in the block-name list	<code>dust</code>

The output slot. Output variable names live in one static array, `var_choice`, and the index of the matching name (`ivar`) is used only by the sanity guards that test numeric ranges (“151–153 are particle outputs, so a `<particles>` block must exist”). Each fork had appended one name at position 153 and written a guard for it. Resolution: `dust_d` keeps 153, `grav_phi` becomes 154, `NOOUTPUT_CHOICES` 154 $\rightarrow$ 155, the gravity guard tests 154. Everything downstream selects by name string, so the number appears nowhere else. The 31 auto-merged files were verified identical to the rehearsal.

4.3 Follow-up 1: task order

Upstream reordered the hydro chain to `Prolongate` $\rightarrow$ `ApplyPhysicalBCs` $\rightarrow$ `ConToPrim` (so physical boundaries act on prolonged ghosts); the dust graph still had the July order. Commit `a7c0def3` aligns `dust/dust_tasks.cpp`:96-98. Inert on uniform meshes (dust is fatal on refined ones), but the invariant of §3.4 says keep them in lockstep.

4.4 Follow-up 2: the deposit exchange on the rank-packed path

Symptom. The first multi-rank test of the merged tree (the damping problem with drifting particles, `v0 = 0.5`, 4 ranks) hung at cycle 0 with all four processes spinning; the 1-rank run of the same binary reproduced the reference number exactly, and multigrid tests at 4 ranks were bit-identical to serial.

Mechanism. `MeshBoundaryValuesDep` (`bvals/bvals_dep.cpp`) was written as a copy of July’s `MeshBoundaryValuesCC`: per-neighbor `MPI_Isend` on `sendbuf[n].vars_req[m]` with `CreateBvals_MPI_Tag(lid, dn)` tags, and per-neighbor `MPI_Test` on `recvbuf[n].vars_req[m]`, relying on the base class `InitRecv` to post the matching per-neighbor receives. In this tree `MeshBoundaryValues::InitRecv` (`bvals/bvals_tasks.cpp`:32) is *rank-packed* since upstream #775: it builds, once per mesh sequence, a metadata table of all (block, neighbor) payloads per peer rank (`BuildRankPackedVarMetadata`, `bvals/bvals.cpp`:134, which also does a one-shot header exchange so receivers know the pack order), and posts *one* `MPI_Irecv` per peer rank with tag 1 into `rank_recvbuf_vars_`. `ClearRecv/ClearSend` wait on those aggregated requests. So the deposit’s per-neighbor sends had no receiver, its tests polled requests that were never posted, and the stage never completed.

Fix. Commit 774c6519 re-expresses the two transport steps in the same idiom as the cell-centered exchange (`PackAndSendCC` and `RecvAndUnpackCC` of `MeshBoundaryValuesCC`):

- `PackAndSendDeposit` (:122): on-rank neighbors are still written straight into the destination's `recvbuf[dn].vars`; off-rank payloads go into `rank_sendbuf_vars_` at `send_agg_offset_ - (m*nngnbr+n)`; after one fence, one `MPI_Isend` per message in `send_var_msgs_` (tag 1, `send_var_reqs_`);
- `RecvAndSumDeposit` (:228): `MPI_Test` over `recv_var_reqs_`; in the (deliberately scalar) loop over neighbors the addend is `aggrbuf(base + bi)` when `recv_agg_offset_ - (m*nngnbr+n) >= 0` and `rbuf[n].vars(m, bi)` otherwise.

The metadata builder sizes each entry from the buffer index ranges the derived class defines (`GetVarDataSize` $\rightarrow$ `isame_ndat`), so the reversed data flow of the additive exchange (pack ghosts, sum into active strips) needs no change to the base class. The serial path is unchanged (bit-identical damping ladder); ranks 1, 2 and 4 now agree to every printed digit.

Messaging. Every new boundary-exchange class in this tree must use the rank-packed vars path (`send_var_msgs_ / recv_var_reqs_ / *_agg_offset_`); the per-neighbor `vars_req` arrays are dead for vars traffic and only the flux-correction path still posts per-neighbor requests.

4.5 The second merge: catching up with the dust fork

What came in. A question about a colleague's commit revealed that the local dust clone was two weeks behind. `origin/dust` at 1f3e1ed9 adds, oldest first: a CPU particle-mesh optimization (plain addition instead of atomics under `Kokkos::Serial`, weights evaluated in cell coordinates); the NSH generator without dust; *their own* import of upstream #775 (0d7688b6), extended to the additive deposit exchange — the same rewrite as 774c6519, validated by them at 2, 4 and 8 ranks and on CUDA; adaptive PC2 / split-BE coupling (4b44ba09); and direct PC2 (7fc41cd8).

Conflicts. Three files. `bvals/bvals_dep.cpp`: both sides rewrote the same functions the same way; the fork's version was taken (it additionally rebuilds the rank-packed metadata inside `PackAndSendDeposit` when the variable count or mesh sequence changed, instead of relying on `InitRecv` having run). `bvals/bvals_fc.cpp`: ours kept — the early return on inactive face-centered faces is upstream #739's $\nabla \cdot B$ fix, which the fork's #775 import predates. `mesh/mesh_refinement.cpp`: ours kept — the fork calls `InitBoundaryValuesAndPrimitives` at the ancestor's position, while the Phase-0c code here calls it later, after the shear-state rebuild; keeping both would call it twice. The base boundary-value files auto-merged with *no net change*: the fork's #775 import is textually the upstream copy this tree already had.

The PC2 and hybrid couplings. `<dust>/coupling` now selects `imex` (the default, §3.3, requires `imex2+`), `hybrid`, or `pc2` (both require `rk2`; `gamma_switch` is ignored with a warning). From the fork's own description: the *PC2* branch advances the particles once per cycle without a coupled gas solve — a fresh old-state drag-feedback predictor is deposited during RK stage 1, the midpoint gas velocity is built after the completed stage-1 boundary tail, the particles take an implicit-midpoint drag update at the predicted midpoint position, the actual drag impulse is deposited at that midpoint, and the final gas momentum is corrected by removing the tableau-weighted predictor. Particles migrate once, after the update. The *split-BE* branch first completes a dust-free hydro RK2 step

and then performs one full- Δt coupled backward-Euler drag solve (first order, L-stable). `hybrid` switches between them globally with hysteresis on two stiffness measures,

$$\zeta = \max_p \frac{\Delta t}{t_{s,p}}, \quad \chi = \max \frac{\Delta t \rho_d}{\rho_g t_s}, \quad (1)$$

entering split-BE when either exceeds `hybrid_enter_{zeta,chi}` (0.5) and returning to PC2 only when both fall below `hybrid_exit_{zeta,chi}` (0.25); `hybrid_force_mode = pc2 | split_be` pins a branch, and `coupling = pc2` is that pin made permanent. The choice is reduced across ranks so every rank takes the same branch, and a `DUST_HYBRID_SUMMARY` line reports the cycle counts and the last ζ , χ . In the task graph (`dust/dust_tasks.cpp`) this appears as per-stage gates on every exchange and two migration chains; the deposited $Q, \mathbf{P}$ field gains a fifth feedback-rate channel; a second timestep task follows the split-BE migration.

Consequence for the program. The science targets ($St = 0.1\text{--}10$ at $\Delta t \sim 10^{-2} \Omega_0^{-1}$) sit in the resolved regime $\zeta \ll 1$, so `coupling = pc2` under `rk2` is the natural production setting: second order, about twice the speed of the IMEX dc1 path by the fork’s measurement, no `imex2+` CFL penalty, and the same integrator as all gravity validation. Note that `hybrid` with the three-species NSH test ($t_{s,\min} = 0.01$) already trips $\zeta = 0.55$ and runs split-BE throughout; the validation document records the first-order signature that produces. The IMEX path stays as the reference for stiff configurations.

4.6 Build notes

`imex2+` needed nothing: it is upstream #612 and was already in the tree; the dust fork only refactored its coefficients into `SetImEx2PlusCoefficients`. Dust sources are part of every build (`src/CMakeLists.txt`); the four dust problem generators are built-in (`-D PROBLEM=built_in_pgens`, selected by `<problem>/pgen_name`). The rehearsal clone was built without network by linking the tree’s `kokkos` checkout and pointing `FETCHCONTENT_SOURCE_DIR_KOKKOS-FFT` at the cached `build/_deps/kokkos-fft-src`.

5 Phase 4b (i): the shear-periodic deposit fold

5.1 The problem

A particle near a `MeshBlock` boundary deposits part of its weight into ghost cells of its own block; the additive exchange (§4.4) adds those ghost deposits into the active cells of the neighbor that owns the location. Across a shear-periodic x_1 face the owner’s cells sit at a *different* y : the tree wraps `shear_periodic` exactly like `periodic` (`mesh/meshblock_tree.cpp`), so the x_1 buffers of the face blocks carry the deposits to the unshifted y — the “exact only for azimuthally uniform states” caveat the module used to print. The copy exchanges of hydro and of u^* fix this by *overwriting* the x_1 ghosts with a remapped copy after the plain exchange; an additive exchange cannot overwrite, so the unsheared contribution has to be suppressed and replaced.

5.2 The sign

The shearing-box identification used throughout the tree is $u(x, y) = u(x + L_x, y - q\Omega_0 t)$: the copy exchange fills the inner- x_1 ghosts with the outer boundary’s content shifted by $+y_{\text{sh}}$ ($y_{\text{sh}} = q\Omega_0 L_x t$), and the outer ghosts by $-y_{\text{sh}}$ (`multigrid/multigrid_shear.cpp:144-146`, `gravity/fft_gravity.cpp:414-416`). A *deposit* in an inner ghost cell at y is a contribution to the outer boundary at $y - y_{\text{sh}}$; the deposit map is the adjoint (transpose) of the ghost-fill map, so its signs are reversed:

Adjoint signs. Inner-ghost deposits are shifted by $-y_{\text{sh}}$ and added to the *outer* face blocks; outer-ghost deposits by $+y_{\text{sh}}$ into the *inner* face blocks. Copying the multigrid fill’s per-face signs verbatim gives the wrong answer.

5.3 The fold

`MeshBoundaryValuesDep::FoldShearDeposit(a, yshear, rcon)` (`bvals/bvals_dep.cpp`) reuses the Phase-1 global-plane machinery of `Multigrid::FillShearGhostPack` with three changes: the gather reads *ghost* slabs and is additive, the signs are the adjoint ones, and the scatter *adds* into active cells.

1. **Gather.** Two planes `shear_plane_(face, v*ng+d, gk, gj)` of the global y - z extent (uniform grid: `mesh_indcs.nx2×nx3`), zeroed, then every face block adds its whole x_1 ghost slab — all k, j including the y/z corner ghosts — at the periodically wrapped global index (face 0 $\leftarrow$ inner ghosts `is-1-d`, face 1 $\leftarrow$ outer ghosts `ie+1+d`). The corner rows overlap the active rows of the y/z neighbors’ slabs, so the gather is an atomic *sum*: unlike the multigrid fill, several blocks legitimately write the same plane cell. Rows beyond a non-periodic z face are dropped (the vertical-boundary policy of §10 item 7; the dust guards currently require periodic z).
2. **MPI seam.** `Kokkos::fence` and one `MPI_Allreduce(SUM)` of each plane, reached by every rank in lockstep (the fold is a task; all ranks run the same graph, and every send a rank waits for is posted before the collective).
3. **Remap.** Per (face, variable, depth, z -row): the integer part of the shift is folded into a periodically wrapped scratch load, the fractional part `eps = fmod(ysh, dy)/dy` goes through `DC/PLM/PPMX_RemapFlx` (`shearing_box/remap_fluxes.hpp`) and the conservative update $q_j \leftarrow q_j - (F_{j+1} - F_j)$; $dy = L_y/\text{gny}$ is the global rank-consistent cell width (the `ConsistentDx2` lesson). Face 0 uses $-y_{\text{sh}}$, face 1 $+y_{\text{sh}}$.
4. **Scatter-add.** Face 0 into the outer face blocks’ `ie-d`, face 1 into the inner blocks’ `is+d`, active k, j only.

The remap preserves the row sum, so the total deposited mass is conserved to round-off; the summation order in the gather and the reduction can differ with the rank count at round-off level only.

5.4 Suppressing the unsheared contribution

`RecvAndSumDeposit` skips, on a block whose inner (outer) x_1 face is shear-periodic, every receive buffer whose direction has $o_{x1} < 0$ (> 0): faces, x_1x_2 and x_3x_1 edges, and corners — all of which carry x_1 -ghost content of the neighbor across the face. The direction of buffer index n comes from a 56-entry table built in the constructor from the slot layout documented in `mesh/nghbr_index.hpp` (faces 0–7, x_1x_2 edges 16–23, x_3x_1 edges 32–39, corners 48–55) and cross-checked against `NeighborIndex` at construction.

A pitfall worth recording. The first version built that table by calling `NeighborIndex` over all direction and sub-index arguments. For edges the function uses only the first sub-index, and the second landed on *other* slots — the x_3x_1 edges of one sign overwrote those of the other. The plain pass then still delivered the unsheared x_3x_1 -corner deposits, which the unit test caught as an off-by-one in the multiplicity check (§5.6) and as a 10^{-4} drift of the 3D NSH equilibrium.

5.5 Hooks and parameters

Two tasks, `DustGasDrag::FoldDepQP` and `FoldPMBR` (`dust/dust_tasks.cpp`), follow `RecvDepQP` and `RecvPMBR` with the same per-coupling stage gates; the implicit solve now depends on the fold. `CompleteAddExchange` (coupled solvers) calls the fold as well. The offset is `DustGasDrag::ShearOffset() = q\Omega_0 L_x t` at `pmesh->time` on every stage, the convention of the u^* remap. `<dust>/shear_remap` selects the order (`plm` default, as for u^* ; `ppmx` available), `<dust>/shear_fold = false` is a diagnostic that skips the fold (the face blocks' x_1 deposits are then dropped, not unsheared). The warning the module printed under shear is replaced by an information line.

5.6 The test generator

`pgen_name = dust_deposit_shear` (`pgen/tests/dust_deposit_shear.cpp`, `inputs/dust/dust_deposit_shear.athinput`) exercises the exchange without particles: the x_1 ghost slabs of the face blocks are filled with known data, every other cell is zero, and the exchange runs exactly as the module runs it. Ghost deposits are *partial* sums (a block's corner ghosts overlap its neighbors' active rows), which the first version of the test got wrong by filling every ghost cell with the full function; the corrected test carries three fields: a smooth $g(x, y, z)$ periodic in y, z on the slab rows inside the block's own y/z range (one source per location, so the strips must equal g at the shifted location up to the remap error), a constant on the same rows (exact for any shift), and a constant on the *whole* slab whose received value must equal the analytic overlap multiplicity $(1 + n_y)(1 + n_z)$ of the source row (exact at integer shifts). The generator also reports the largest value leaked into any active cell outside the receiving strips (must be exactly zero) and the relative mass change (round-off). `problem/fold = false` runs the plain pass alone.

5.7 A structured, rank-independent initial condition

The NSH generator's random placement hashes each particle's tag to a block index *within the local pack*, so the realization depends on the rank count (as does `dust_damping`'s, seeded by the pack's first GID) — fine for physics runs, useless for serial-vs-MPI comparisons; a divergence that this produced was first mistaken for a bug in the fold (the no-back-reaction control agreed to 10^{-13} because test particles never feed positions back). `dust_nsh` therefore gained `<problem>/mass_mod_amp`, `mass_mod_phase`: on the lattice, particle masses are multiplied by $1 + A \sin(2\pi(y - y_{\min})/L_y + \phi)$, a deterministic azimuthally structured dust distribution (with `check_equilibrium = false`). `inputs/dust/dust_nsh_shear3d.athinput` is the tracked 3D shear-periodic NSH input with all Phase-4b keys present for command-line overrides.

6 Phase 4b (ii): the ideal-gas energy update

6.1 The problem

The dust fork applied the drag back-reaction to the gas momentum only and therefore required an isothermal EOS (a fatal guard). Gravito-turbulence needs cooling, hence an ideal gas, and then every change of the gas momentum must be accompanied by a change of the total energy: otherwise the internal energy $e = E - |\mathbf{M}|^2/2\rho$ silently absorbs the negative of the kinetic-energy change.

6.2 The rule

`GasKick(u0, m,k,j,i, dm, ideal)` (`dust/dust.hpp`) is the single entry point for a drag kick $\Delta \mathbf{M}$ on a cell: for an ideal gas it adds $\Delta E = (|\mathbf{M} + \Delta \mathbf{M}|^2 - |\mathbf{M}|^2)/2\rho$ before adding $\Delta \mathbf{M}$, so the kick leaves e unchanged. The four sites where the module changes the gas momentum all use it: the IMEX back-reaction apply (`ApplyPMBR`), the IMEX history term (`AddDragHistoryGas`, $\Delta \mathbf{M} = \tilde{a}_{22} \Delta t \mathbf{R}_g$), the PC2 predictor apply (`GasImplicitSolve`, stage 1) and the hybrid commit (`ApplyPMBR`, stage 2, with the tableau-weighted predictor removal folded into the same kick). The coupled solvers only read the gas state and apply through `ApplyPMBR`.

Frictional heating (`<dust>/drag_heating`, default false). The dissipation of one drag kick $\Delta \mathbf{v}_j = c_j(\tilde{\mathbf{u}}_j - \mathbf{v}_j)$ on particle j — the kinetic energy lost by particle and gas together beyond the gas kinetic-energy change that `GasKick` books — is

$$Q_j = m_j c_j (1 - c_j/2) |\tilde{\mathbf{u}}_j - \mathbf{v}_j|^2 \geq 0, \quad (2)$$

deposited with the particle’s weights into a fourth channel of the back-reaction field `dmom` (now four variables; the exchange objects follow) and added to E where the momentum channels are applied; the IMEX history term scales it with the same $\tilde{a}_{22}$. On the PC2 path the midpoint impulse gives $Q_j = m_j [\Delta \mathbf{v}_{\text{drag}} \cdot (\mathbf{u}_g - \mathbf{v}_h) - \frac{1}{2} |\Delta \mathbf{v}_{\text{drag}}|^2]$, the midpoint quadrature of $m_j |\mathbf{u} - \mathbf{v}|^2/t_s$ over Δt in kinetic-energy form. With an isothermal EOS the flag is ignored with a warning.

6.3 What the bookkeeping cannot do, and the lesson from the first test

Inside a low-storage RK stage the conserved state is combined linearly, $\mathbf{u} \leftarrow \gamma_0 \mathbf{u} + \gamma_1 \mathbf{u}_1$; E combines linearly but $|\mathbf{M}|^2/2\rho$ does not, so the combined internal energy exceeds the combination of internal energies by $\gamma_0 \gamma_1 |\mathbf{M}_a - \mathbf{M}_b|^2/2\rho$ per step, where $\mathbf{M}_a - \mathbf{M}_b$ is the stage-1 momentum change — for a drag-dominated cell, the stage-1 kick. The default mode therefore does *not* keep e exactly fixed over a step; it gains this positive gap, first order in Δt (the kick scales with Δt while $\Delta t \ll t_s$) and sub-first-order while $\Delta t \gtrsim t_s$ of the stiffest species, where the kick saturates. The particle velocities are combined the same way, so the heating mode carries the analogous gap on the dust side: exact conservation of $E_{\text{gas}} + K_{\text{dust}}$ is not attainable in this framework, and the measured residuals (validation document §6) are this truncation error. Both modes converge with Δt ; both are inert for an isothermal gas.

The first energy test looked like a 30% bookkeeping error; it was the history file’s six significant digits (a total energy of 2.5 cannot show changes of 10^{-6}). `inputs/dust/dust_damping_ideal.athinput` therefore uses $p_{\text{gas}} = 10^{-5}$, comparable to the dissipated energy, and `<time>/dtmax` to pin the step. The NSH generator gained the ideal-gas initialization it lacked (uniform `<problem>/pgas`); `inputs/dust/dust_nsh_shear3d_ideal.athinput` is the ideal-gas twin of the 3D input.

7 Phase 4b (iii): the particle timestep under shear

7.1 The problem

The inherited transport limit (§3.6) is $\Delta t \leq c_p \min_p \Delta x_2/|v_y - q\Omega_0 x|$ with $c_p = \text{dt_cfl} = 0.5$: one cell of the *total* azimuthal motion per step. The background shear is a smooth translation the drift integrates exactly ($y \leftarrow y + \Delta t (v_y - q\Omega_0 x)$ is linear in x), not a signal the mesh has to resolve; the gas, orbitally advected, is limited by its peculiar velocity only. In the $L_x = 1$ test boxes the two limits happen to coincide with the sound-speed limit, which is why the penalty never showed. In an

$L_x = 8$ box with the NSH input's $\Delta x_2 = 1/32$ the particle limit is $0.5 \cdot (1/32)/(1.5 \cdot 4) = 2.6 \times 10^{-3}$ against a gas limit of 1.37×10^{-2} : a factor 5.3, growing with L_x .

7.2 What the limit is actually for

Two mechanisms depend on how far a particle moves in one stage, and neither is a cell:

1. **The deposit halo.** The TSC stencil reaches one cell beyond the particle, and `MeshBoundaryValuesDep` exchanges one ghost layer. Both are satisfied whenever the particle lies *inside its owning block* at deposit time — which is what the per-stage migration establishes: after `ParticlesBoundaryValues::SetNewPrtclGID` every particle carries the `gid` of the block that contains it. The halo-violation abort in `DepositDrag` (`dust/dust_pm.cpp`) is precisely the detector of a failed migration.
2. **The migration.** `SetNewPrtclGID` forms the integer offset of the particle from its block, $i_x = \lfloor (x - x_{\min} + L_x^{\text{mb}})/L_x^{\text{mb}} \rfloor - 1$ and likewise in y, z , and indexes the neighbor table with `NeighborIndex( $i_x, i_y, i_z, \dots$ )`. The table covers faces, edges and corners: offsets in $\{-1, 0, 1\}$. A particle that traverses a whole `MeshBlock` in one stage produces $|i_y| = 2$, an out-of-range neighbor slot, and a wrong or invalid destination. The one exception is a crossing of the shear-periodic x_1 mesh boundary, routed by the (y, z) shear maps of Phase 4a for any azimuthal offset.

So the hard requirement is *less than one MeshBlock per stage*, in every direction, on the total transport velocity; the cell-per-step limit on the peculiar velocity stays for accuracy, as for the gas.

7.3 The rule

`DustGasDrag::ComputeNewTimeStep` (`dust/dust_newdt.cpp`) now reduces two minima over the particles in one kernel and takes the smaller:

$$\Delta t_{\text{cell}} = c_p \min_p \min_i \frac{\Delta x_i}{|v_i^{\text{cell}}|}, \quad v_y^{\text{cell}} = \begin{cases} v_y & \text{dt_transport} = \text{relative} \\ v_y - q\Omega_0 x & \text{full} \end{cases} \quad (3)$$

$$\Delta t_{\text{block}} = \frac{s}{\beta_{\max}} \min_p \min_i \frac{L_i^{\text{mb}}}{|v_i^{\text{full}}|}, \quad \beta_{\max} = \max_{\text{stages}} |\beta_s|, \quad (4)$$

with $v_y^{\text{full}} = v_y - q\Omega_0 x$ in the 3D shearing box (and $v^{\text{full}} = v$ otherwise), L_i^{mb} the block extent, and $s = \text{<dust>/dt_block_safety}$ (default 0.5). The factor $\beta_{\max}$ is what makes s the literal fraction of a block crossed per stage. In a low-storage stage the position is combined as $y \leftarrow \gamma_0 y + \gamma_1 y_1 + \beta_s \Delta t \dot{y}$ with $\gamma_0 + \gamma_1 = 1$, so the displacement from the position the stage started with is $\gamma_1(y_1 - y) + \beta_s \Delta t \dot{y}$; for the tableaux in use the register term never adds to the drift term in magnitude, and the per-stage displacement is bounded by $\beta_{\max} \Delta t |\dot{y}|$:

integrator	stage displacements / ($\Delta t \dot{y} $)	$\beta_{\max}$	bound used
<code>imex2+</code> , $\gamma = 1 + 1/\sqrt{2}$	1.707, 0.707	1.707	1.707
<code>imex2+</code> , $\gamma = 1/2$ (switch)	0.5, 0.5	1.0	1.0
<code>rk2</code>	1.0, 0	1.0	1.0
<code>rk3</code>	1.0, 0.5, 0.5	1.0	1.0

The kernel reads `pdrive->beta[]` over `nexp_stages`, so the γ -switch is covered without a table. The premise $\gamma_0 + \gamma_1 = 1$ with non-negative weights holds for every tableau the module’s integrator guard admits (`imex2+` for `coupling = imex`, `rk2` for `pc2/hybrid`) and for `rk3`; `rk4` has $\gamma_0 < 0$ and would need its own bound, but cannot be selected. Note the pre-existing asymmetry the cell limit carries under `imex2+`: $c_p = 0.5$ is applied to Δt , while stage 1 advances $\gamma\Delta t = 1.71\Delta t$, i.e. 0.85 cells per stage; it is unchanged here and harmless (nothing in the module needs less than a block).

7.4 Parameters and diagnostics

`<dust>/dt_transport = relative|full` (default `relative`; `full` is the legacy limit, bit-identical to Phase 4b (ii)) and `<dust>/dt_block_safety` (default 0.5; a value ≥ 1 is accepted with a warning for the diagnostic run of validation §7 that shows the halo abort the guard prevents; the constructor prints which limit is in force in a 3D shearing box). The guard is always on, in all three directions and in every geometry; outside a shearing box it is inert in practice (the cell limit is stricter by N_i^{mb}). The destructor prints `DUST_DT_SUMMARY block_guard_cycles=n`, the number of cycles in which the block guard was below the cell limit *on rank 0* — a rank-local count that says which of the two particle limits was the active one, not whether the particles limited the run at all (the gas usually does once the shear is out). The two 3D shear-periodic inputs carry both keys so they can be varied from the command line (§11).

7.5 What the larger step exposes

Once the shear is out of the limit, the step in a wide box is the gas step, and for species with t_s below it the IMEX coupling runs in its stiff regime. The validation scan (§7 there) shows this is a plain, converging truncation error — second order, 1% in the quiet phase and 26% during a streaming burst at $\Delta t/t_{s,\text{min}} = 1.4$ — that the legacy limit had been masking in wide boxes only. The operating rule that follows is recorded in the validation document; it does not change the timestep code, which now limits what must be limited and nothing else.

7.6 What does not change

The push, the deposits, the fold and the migration are untouched: a larger step simply means that particles now hop MeshBlocks routinely in y (up to $s N_2^{\text{mb}}$ cells per stage) and cross the shear-periodic face with multi-cell azimuthal offsets, both paths that existed and were exercised before at one cell per step. The split-BE branch of the hybrid coupling keeps its later timestep task (`NewTimeStep2`), which calls the same routine.

8 Phase 4c: dust self-gravity

8.1 Design, as built

The design of §2 is implemented without change of plan: one Poisson solve per stage on the *total* density, the gas feeling the total potential through its existing source term, the particles through a mesh force gathered with the deposit kernel. The pieces, in the order they run each active stage (figure of §2.2):

1. **Deposit** (`DustGasDrag::DepositMass`, `dust/dust_gravity.cpp`), in the `before_stagen` list: $\rho_d = \sum_p m_p W(\mathbf{x}_p)/V$ with the module’s PM kernel, at the pre-stage positions — the same

time level as the gas density u_0 the solve reads. Particles sit inside their owning blocks (last stage’s migration), so the stencil reaches at most one ghost layer.

2. **Additive exchange and fold:** `MeshBoundaryValuesDep` (one variable) sums the ghost deposits into the neighbours and `FoldShearDeposit` takes them across the shear-periodic faces exactly as for the drag deposits (§5).
3. **Ghost fill:** a copy exchange (`MeshBoundaryValuesCC`, one variable) with the shear remap (`ShearingBoxCC`) fills the ghost layers of ρ_d from the neighbours’ active cells, because `Multigrid::LoadSource` loads the active zone *plus one ghost layer* of its source (`multigrid.cpp:290`), which for the gas is valid by the hydro boundary fill. The shear offset uses the module’s `ShearOffset() = q\Omega_0 L_x t`, the same instant the multigrid freezes its own phase at (`mg_gravity.cpp:325`).
4. **Registration:** `Gravity::RegisterExtraDensity( $\rho_d$ )` (once, in the dust constructor; `pgrav` is built before `pdust`). `Gravity::Solve` then forms $\rho_{\text{tot}} = u_0(\text{IDN}) + \rho_d$ over every cell into its own array, and the three readers of §2.2 — `LoadSource`, the slab-plane gather (`ComputeSlabPlanes`, which gained an index argument) and the FFT `GatherDensity` — read `SourceArray() (m, SourceIndex(), k, j, i)`: the gas array at IDN when nothing is registered, the total at index 0 otherwise. Neither solver has any dust-specific code.
5. **The solve**, called by the driver between the two lists, unchanged — except that with the dust module present it is *skipped on the dormant imex2+ assembly stage (driver.cpp)*: the module deposits nothing there, so the density would be stale, and nothing consumes ϕ ($\beta_3 = 0$). Gas-only runs keep their stage-3 solve, so their histories are unchanged.
6. **Force** (`ComputeForceField`), first in the `stagen` list: $\mathbf{g} = -\nabla\phi$ by centred differences on the active zone (ϕ is valid on the active zone plus one ghost layer for both solvers), with a zero-gradient fill of the ghost layers at physical (non-periodic) faces; then a copy exchange (three variables) and the shear remap, the pattern of the module’s u^* field.
7. **Gather**, in `ExplicitPush`: the IMEX branch gathers $\mathbf{g}$ at the pre-stage position with the deposit kernel and adds it to the explicit force next to the Coriolis terms; the PC2 branch gathers it at the midpoint position with the weights it already has for u^* and adds $h\mathbf{g}$ to the right-hand side of its implicit-midpoint velocity solve, as it does the vertical tide.

The gas needs nothing new: `<hydro_srcterms>/self_gravity` applies $-\rho\nabla\phi$ with the same centred gradient. The drag back-reaction deposits only the drag kick, so there is no double count of gravity.

8.2 Momentum conservation, exact by construction

With one kernel W for scatter and gather, the total force on the dust is $\sum_p m_p \sum_c W_{pc} \mathbf{g}_c = \sum_c \rho_{d,c} V \mathbf{g}_c$, and the total on the gas is $\sum_c \rho_{g,c} V \mathbf{g}_c$; their sum is $\sum_c \rho_{\text{tot},c} V \mathbf{g}_c$. In the periodic solve $\rho_{\text{tot}} \propto L\phi$ with L the symmetric seven-point Laplacian and $\mathbf{g}$ the antisymmetric centred difference, so $\sum_i (L\phi)_i (\phi_{i+1} - \phi_{i-1}) = 0$ identically: the coupling exchanges momentum between gas and dust without creating any, to the tolerance the solver reaches. The validation document measures 10^{-13} (§8 there) for NGP and TSC alike. This is the reason for insisting on one kernel and on the mesh gradient rather than a direct particle–particle or particle–potential interpolation.

8.3 Synchronous entry points

Problem generators that call the solve outside the time loop use `AssembleGravitySourceNow()` and `ComputeGravForceNow()`, which run the same deposit, exchanges, fold and force with spin-waited receives (the idiom of the coupled drag solvers' `Complete*Exchange`). The static twin tests are built on them.

8.4 Parameters, guards, boundaries

`<dust>/gravity` (default: true when a `<gravity>` block exists), `gravity_source` (the dust in the Poisson source) and `gravity_force` (the particles feel ϕ); test particles in a gas potential are `gravity_source = false`. Dust gravity needs a 3D mesh. The boundary guard of the module now admits physical (non-periodic) x_3 faces in 3D with a warning: ghost deposits beyond such a face have no receiver and are discarded, the density ghosts there are zero and the force ghosts zero-gradient, and a particle that leaves through the face is not routed — the deposit halo check aborts on it. The vertical policy (removal or reflection, and the mass account) remains the Phase-4d item it was.

8.5 What was found on the way

Particles::dtnew uninitialized. The upstream defect of the survey (§10, last item) bit: the constructor never sets `dtnew`, the dust generators each set it by hand, a new generator that did not saw `Mesh::NewTimeStep` read garbage before the first cycle, and the $2\times$ growth rule then held $\Delta t = 0$ forever. Fixed at the source: the constructor initializes it to the float maximum, and `Driver::Initialize` computes the dust transport limit from the initial state as it does for the fluids. The existing runs are bit-identical.

Multigrid geometry. The multigrid requires *cubic cells* ($\Delta x = \Delta y = \Delta z$: its smoother is the isotropic seven-point operator) and, with periodic directions, *more than one MeshBlock across each* (a block that is its own neighbour breaks the level-boundary exchange); the slab-open sheet test gives errors of 10^2 – 10^3 otherwise and 4×10^{-11} on $32 \times 32 \times 64$ cells in a $0.5 \times 0.5 \times 1$ box. The FFT solver has neither restriction. Both Phase-4c dynamic inputs are shaped accordingly; the requirement belongs in the multigrid document.

A numerical instability of the drag coupling. The 3D NSH lattice, exact to round-off for ten orbits at $32 \times 32 \times 8$ (Phase 4b), grows a round-off perturbation by $1.6\times$ per step at $32 \times 32 \times 32$ ($\Delta z = \Delta x$) and at $32 \times 32 \times 24$, but not at $32 \times 32 \times 16$ or $32 \times 32 \times 64$ (validation §8, table). The scan localizes it: it needs the drift ($\eta v_K \neq 0$), the TSC kernel (CIC and NGP are stable at the same step) and $c_s \Delta t / \Delta z \gtrsim 0.35$ evaluated with the *cycle* step — `imex2+` at `cfl_number` 0.35 is stable and 0.40 is not, `rk2/PC2` at 0.30 is stable and 0.45 is not, so it is not the stage Courant number of the gas — and it does not need the stiff species. The growth rate ($\sim 35\Omega$) rules out the streaming instability, which is also absent with CIC. It is a property of the inherited coupling, not of the gravity, and Phase 4b never met it because every 3D box so far had $\Delta z \geq 2\Delta x$. The operating rule until it is understood: `cfl_number` ≤ 0.3 in boxes with $\Delta z \leq \Delta x$ (the gravity NSH input carries it); the mechanism — an aliasing of the three-cell deposit with the vertical sound crossing is the natural suspect — is a Phase-4d item.

9 Phase 4d: the instability identified, the vertical policy, restart, diagnostics

9.1 The instability of §8 is the hydro scheme’s

Two diagnostics settle what the scan of §8 could not. First, the eigenmode: binary outputs of the unstable 32³ dust run (validation §9) show the growing perturbation in the *gas* velocity and density at the grid Nyquist in all three directions, $(k_x, k_y, k_z) \approx (16, \pm 15, \pm 15)$ of 32 — a 3D checkerboard — while the NGP-binned dust density stays flat until particles cross cells. Re-read with that mode’s frequency, $\omega = 2c_s \sqrt{\sum_i \Delta x_i^{-2}}$, every row of the scan falls on one line, $\omega \Delta t \gtrsim 1.2\text{--}1.3$, including the Δz dependence. Second, the decisive test: the same box with *no particles*, its gas velocities seeded with 10^{-10} white noise (`<problem>/gas_vpert` of the NSH generator), grows the same mode at `cfl` 0.45 and 0.40 under both `imex2+` and `rk2`, and damps it at 0.30 and at $n_z = 8$. A uniform state has no perturbation to grow, whatever the scheme’s stability — which is why the gas-only NSH “held to round-off” and the dust run, whose deposit round-off seeds every mode, did not. TSC seeded it where CIC and NGP happened not to; the drift made the seed non-symmetric. The instability is the shearing-box hydro’s (HLLC, PLM, orbital advection) 3D checkerboard Courant limit on cubic cells, and the operating rule `cfl_number` ≤ 0.3 in 3D is a hydro rule the SC14 gravito-turbulence runs already obey. The Phase-4c wording “TSC drag-coupling instability” is superseded.

9.2 Vertical policy: outflow removal

The Athena-C reference (`par_strat3d.c`, `zbc_out = 1`) treats the vertical faces as outflow for particles: a particle that crosses one is gone. Implemented in three parts. `ParticlesBoundaryValues::SetNewPrtclGID` (`bvals_part.cpp`) tests, before the neighbour lookup, whether a crossing particle has left the mesh through a *physical* face — neither periodic, shear-periodic nor internal, from `mesh_bcs` — and marks it dead with `PGID = -1` instead of routing it. `Particles::RemoveDead` (`particles.cpp`) compacts the arrays (an exclusive scan over the survivors into fresh arrays) and refreshes the Mesh’s per-rank and total particle counts; it is collective under MPI (an `Allgather`), so every rank calls it whenever a migration ran. The dust tasks `RemoveEscaped` / `RemoveEscaped2` follow the two migrations’ receives in the task graph, and the drag deposit and the back-reaction commit now depend on them; they sum the removed mass first, and the destructor prints `DUST_ESCAPE_SUMMARY removed=n mass=m` (reduced over ranks). Ghost deposits beyond a physical face are discarded as before, and the density and force ghosts there are as in §8. With every face periodic the marking never fires and the existing runs are bit-identical.

9.3 Particle restart

The restart file gains two pieces. In the writer’s serial header (step 3, after the turbulence RNG state) the root writes the per-rank particle counts and the two array widths; their $(n_{\text{ranks}} + 2)$ integers are added to the header size the parallel offsets are computed from. After all MeshBlock data (a new step 5) each rank writes its particle arrays at an offset that follows the lower ranks’ blocks (shared file) or its own MeshBlock data (one file per rank), every value as a `Real` in variable-major order — the integer properties are small, and a device-independent layout matters more than the bytes. The reader mirrors both: the counts reset the Mesh’s bookkeeping and the array sizes before the MeshBlock reads (the constructor had sized them from `<particles>/ppc`), the arrays are read after them, and the problem generator is then called with `restart = true` as before. The dust generators already return early on restart after enrolling their history functions; the stopping-time contract is

re-checked at the first `GammaSwitch` from the restored IPTS.

9.4 Diagnostics

`dust_dpm` is a new output variable (slot 155): the dust density as the module deposits it — its kernel, the additive exchange and the shear fold — obtained from `AssembleDustDensityNow()` (the synchronous assembly of §8, now available in every dust run: `rho_dust` and its exchange objects are allocated unconditionally). `dust_d` remains the NGP binning of upstream. The SC14 history (`GravitoTurbHistory`) gains seven columns when a dust module exists: the dust mass, the three momenta (shear-relative velocities, as the particles carry them), the kinetic energy, the Reynolds stress $\sum mv_x v_y$ and the cumulative escaped mass; the gravitational-stress columns already measure the total potential. The NSH history’s mass-weighted momenta are from Phase 4c.

9.5 The science setup: Baehr, Zhu & Yang (2022) in two stages

The target runs (their Table 1, models `BR_L_10` and `noBR_L_10`) are built in SC14 code units with $H_g = c_{s0}/\Omega = 1.02 \times 2.12625 = 2.16878$ — $Q_0 = 1.02$ through c_{s0} — so their $512^2 \times 256$ box with $L_x = L_y = (80/\pi)H_g$, $L_z = (40/\pi)H_g$ is $55.23^2 \times 27.61$ at $\Delta = 0.1079$ (cubic, $0.05 H_g$). Two additions serve it. `<hydro_srcterms>/bcool_cs2_floor` gives β -cooling an irradiation floor: only the internal energy above the floor temperature $c_s^2 = \text{bcool_cs2_floor}$ is cooled (their background irradiation that holds Q above Q_0 ; zero recovers SC14). And the runs are staged: the gas alone to saturation (`inputs/shearing_box/gravito_turb_baehr.athinput`, $t = 50$, dumps every 10), then a restart that *inserts* the particles — `athena -r dump -i gravito_turb_baehr_dust.athinput`: the add-on’s blocks merge over the dump’s input, `<particles>/restart_insert = true` makes the restart reader skip the particle section the dump does not have, and `ProblemGenerator::GravitoTurbInsertDust` places the layer: `ppc` particles per cell in total (1.5×10^6), uniform in (x, y) , Gaussian in z with width `dust_hz` = H_g (their initial gas width), at rest in the shearing frame, one species at $t_s = 1/\Omega$ ($St = 1$), equal masses summing to $Z = 0.01$ of the gas mass in the box. Per-block counts follow the Gaussian mass in each block’s z range and positions come from a hash of the global block id and the particle’s index, so the insertion is decomposition-independent; the Mesh counts and tags are rebuilt. PC2 coupling under the gas run’s `rk2`, `cf1` 0.3, TSC, self-gravity on, diode faces with the outflow removal. Back-reaction on or off is one command-line switch (`dust/back_reaction`). Continuing a stage-2 run from one of its own dumps needs `particles/restart_insert=false`, since those dumps carry the particles. Job scripts: `scripts/cluster/gt_baehr_stage{1,2}.slurm`, README §3.

10 What the next phases must add

The survey established these gaps; they are listed in the order they bite.

1. **Shear-periodic additive deposit remap (4b). Done, §5.** The Phase-4c gravity deposit will be a third additive field and inherits the fold through `MeshBoundaryValuesDep`. The alternative design (image particles deposited directly in the neighbor’s frame) stays in reserve; the field remap is exact for constants and second/third order otherwise, and it costs one small collective per fold.
2. **Ideal-gas back-reaction (4b). Done, §6.**
3. **Particle timestep under shear (4b). Done, §7:** limit on v_y , with a `MeshBlock` guard on the total velocity.

4. **Dust self-gravity (4c).** **Done**, §8, both solvers. The instability it exposed is the hydro scheme’s 3D checkerboard mode (§9); $cf1 \leq 0.3$ in 3D.
5. **Particle restart (4d).** **Done**, §9.
6. **Diagnostics (4d).** **Done** in part, §9: `dust_dpm` and the SC14 dust history columns. Still open: `pvtk` omits stopping time and mass; a Roche-density / Hill-radius clump finder does not exist.
7. **Vertical boundary policy (4c/4d).** **Done**, §9: outflow removal with the mass accounted.
8. **Refinement (Phase 5).** Dust is fatal on `multilevel` in three places, and upstream particles have no remesh handling at all (`PGID` goes stale after the first regrid; nothing in `mesh_refinement.cpp` touches particles). Needs the deposit exchange across levels and a particle redistribution hook next to `ShearingBox::ReinitAfterMeshUpdate`.
9. **Upstream defect.** `Particles::dtnew` was never assigned by the constructor; **fixed** in Phase 4c (§8).

11 Input-file interface

Required: `<time>/integrator = imex2+` for `coupling = imex, rk2` for `pc2` or `hybrid`; `<particles>/particle_type = dust, pusher = imex_dust, ppc` (particles per cell, an integer multiple of `nspecies` for lattice placement); `<mesh>/nghost >= 2`; periodic boundaries, or shear-periodic x_1 in 3D. Any EOS (the isothermal requirement for back-reaction is gone, §6). Shearing-box runs read `<shearing_box>/qshear, omega0, stratified`.

<dust> parameter	default	meaning
coupling	imex	imex (needs imex2+) / hybrid / pc2 (both need rk2); see §4.5
shear_remap	plm	dc / plm / ppmx: y -remap order of the shear-periodic deposit fold (§5)
shear_fold	true	false: diagnostic, skip the fold
drag_heating	false	ideal gas: deposit the frictional dissipation of each drag kick as heat (§6); default keeps e fixed and discards it
hybrid_force_mode	auto	auto / pc2 / split_be (hybrid only)
hybrid_enter_zeta,	0.5, 0.5	split-BE entry thresholds on ζ , χ
hybrid_enter_chi		
hybrid_exit_zeta,	0.25, 0.25	PC2 re-entry thresholds
hybrid_exit_chi		
nspecies	1	number of species; $\tau_{\text{aus}}1..N$ (positive) required
dust_to_gas	0.01	total dust-to-gas ratio for the default mass initializer
back_reaction	true	deposit the drag reaction onto the gas
deposit	tsc	ngp / cic / tsc assignment and gather kernel
dt_cfl	0.5	particle transport CFL factor (cell-crossing limit)
dt_transport	relative	relative: the cell limit uses the peculiar v_y ; full: the legacy $v_y - q\Omega_0 x$ (§7)
dt_block_safety	0.5	fraction of a MeshBlock a particle may cross per stage on the total velocity (≥ 1 : guard off, warns; diagnostic)
gravity	(<gravity> present)	dust takes part in self-gravity (3D only)
gravity_source	true	the PM dust density joins the Poisson source
gravity_force	true	particles feel $-\nabla\phi$ of the total potential
gamma_switch	false	Krapp et al. eq. 18: $\gamma \rightarrow 1/2$ when $\Delta t > \max t_s$
stopping_time_mode	species_fixed	or particle_static, dynamic
drag_solver	local	local / dc1 / dc2 / pcg / adaptive / applya (coupled Schur-complement solvers; local is the BLKP19 closed form and the validated baseline)
drag_rtol, drag_atol,	10^{-11} , 10^{-14} ,	PCG tolerances
drag_iter_max	200	
drag_adaptive_*	—	error target of the adaptive solver
drag_diagnostic_-interval	0	DUST_SOLVER_DIAG print cadence

Problem-generator parameters (ρ_0 , etavk , eps_s , placement, seeds, eigenmode amplitudes; dust_nsh 's mass_mod_amp , mass_mod_phase , check_equilibrium ; $\text{dust_deposit_shear}$'s time0 , remap , fold) are specific to the five built-in generators dust_damping , dust_nsh , dusty_wave , streaming_linear , $\text{dust_deposit_shear}$; the files under `inputs/dust/` and `../athenak_dust_tests/inputs/` are the reference configurations. Outputs: `file_type = pvtk` with `variable = prtcl_all` (positions, gid, tag, species, velocities); `variable = dust_d` on any mesh output (NGP mass density, needs a <dust> block); <time>/dtmax caps the step for fixed-step convergence tests.