

Dust Particles and Gas+Dust Self-Gravity
in the AthenaK Shearing Box:
I. The Dust-Module Merge (Phase 4a),
the Shear-Periodic Deposit Fold and the Ideal-Gas
Energy Update (Phase 4b, i–ii)

athenak-multigrid tree, branch dust-multigrid

September 3, 2026

Abstract

This document opens the *dust track* of the self-gravity program: Lagrangian dust particles in the shearing-box code with multigrid self-gravity on refined meshes, toward the gas+dust gravito-turbulent boxes of Baehr, Zhu & Yang (2022) with the radial pressure gradient they omitted. It records what the code does on test problems, phase by phase; the mechanisms are in the companion `dust_selfgravity_implementation.pdf`.

Phase 4a merges the sibling fork `athenak-dust` — particle-mesh dust with Benítez-Llambay–Krapp–Pessah implicit drag inside the Krapp et al. (2024) IMEX integrator, validated in its own repository — into the multigrid tree, and asks two questions: is the dust module unchanged by the merge, and is the gravity machinery unchanged by the dust module? Both answers are yes to every printed digit. On the merged tree the dust fork’s own acceptance numbers reproduce exactly: the N -species damping ladder converges at second order in Δt with the reference errors 1.3529×10^{-7} , 3.4627×10^{-8} , 8.7724×10^{-9} , 2.2075×10^{-9} and total momentum conserved to 10^{-16} ; the multi-species Nakagawa–Sekiya–Hayashi drift equilibrium holds to 10^{-16} over ten orbits; the damped dusty sound wave converges toward second order in resolution; the ion-neutral C-shock under `imex2+` is bit-identical to its baseline. The gravity solver’s documented statics reproduce exactly (slab sheet error 1.203×10^{-12} , Tomida & Stone contraction 0.125, boundary-annuli shearing wave 5.950×10^{-3} / 2.057×10^{-3}), bit-identically at 4 ranks. The one defect found — a multi-rank deadlock of the additive deposit exchange against the tree’s rank-packed messaging — is fixed and ranks 1, 2 and 4 agree to all digits. Finally the gas-only stratified gravito-turbulence box evolved with `rk2` and with `imex2+` (the integrator the dust module requires) agrees in every history column to between 2×10^{-6} and 5×10^{-3} at $t = 3 \Omega^{-1}$, the expected spread between two second-order integrators.

A **second merge** the same day brought in the dust fork’s latest seven commits. The whole battery reproduces on the result to every digit (the gas-only run bit for bit), and the fork’s new predictor–corrector coupling under plain `rk2` (`coupling = pc2`) is exercised for the first time here: second order in the damping ladder (orders 2.23, 2.10, 2.04) and in the dusty wave (1.85, 1.98, 2.00), NSH to round-off, rank-independent at 1, 2 and 4 ranks — the integrator freedom that lets dust run under the same `rk2` as every gravity result before it.

Phase 4b (i) adds the conservative azimuthal remap of ghost deposits across the shear-periodic radial faces, the first of the gaps the dust module carried. A particle-free unit test certifies the new fold: the plain exchange delivers nothing across the faces, the fold is exact at integer shifts (3.5×10^{-16}) including the overlap multiplicity of corner rows, mass is conserved to 10^{-15} , nothing leaks outside the receiving strips, the remap error converges at second order

(donor cell), second/third (PLM), third (PPMX) in a fixed-phase scan, and 1, 2 and 4 ranks agree to every digit. The 3D shear-periodic NSH equilibrium still holds to round-off; a deterministic azimuthally structured dust distribution with back-reaction agrees between serial and 4-rank runs to 10^{-14} over an orbit, and shows the old unsheared exchange was a 1–7% error in the gas kinetic energy while the remap order matters only at 10^{-7} .

Phase 4b (ii) removes the isothermal restriction on back-reaction. With every drag kick booking its kinetic-energy change, the ideal-gas damping test keeps the momentum dynamics identical to the isothermal reference and changes the internal energy by 1% (IMEX) / 2% (PC2) of the dissipated energy at a step equal to the stiffest stopping time — the low-storage RK combination gap, converging with the step; the frictional-heating option conserves $E_{\text{gas}} + K_{\text{dust}}$ to 4–6% of the dissipated energy at that step, first order in Δt ; and on the 3D shear-periodic NSH state the heating rate matches the analytic dissipation to 1.5%, 0.7%, 0.4% as the step is halved twice, with the equilibrium itself at round-off in both modes.

Contents

1	Plan and scope	3
1.1	The program	3
1.2	What Phase 4a must establish	3
2	Test setup	3
2.1	Builds	3
2.2	Inputs and commands	4
3	Results	4
3.1	Dust fidelity	4
3.2	Regressions	6
3.3	Rank independence, and the deadlock	6
3.4	Gravity regression	6
3.5	The integrator twin	7
4	The second merge: regression, and the PC2/hybrid couplings	7
4.1	Regression	7
4.2	PC2 coupling under <code>rk2</code>	7
4.3	Hybrid coupling and the split-BE branch	9
5	Phase 4b (i): the shear-periodic deposit fold	9
5.1	The unit test	10
5.2	Regression: the 3D shear-periodic NSH equilibrium	10
5.3	A structured run: rank independence and the size of the correction	11
6	Phase 4b (ii): the ideal-gas energy update	11
6.1	The damping test with an ideal gas	12
6.2	The NSH state with an ideal gas	12
6.3	Regression	13
7	Summary	13

1 Plan and scope

1.1 The program

The science target is the interplay of gas gravitational instability and dust concentration: 3D stratified, cooling, gravito-turbulent shearing boxes with superparticles (Baehr, Zhu & Yang 2022: $512^2 \times 256$, $\beta = 10$ cooling with an irradiation floor, $\gamma = 5/3$, $\text{St} = 0.1, 1, 10$, $Z = 0.01$, $Q_0 = 1.02$), extended by the radial pressure gradient that admits the streaming instability, and eventually adaptive zoom on the dust clumps — the configuration their limitations section names. The route (implementation document, §1.2): Phase 4a merges the dust module; 4b adds what the 3D shearing box needs from it (shear-remapped additive deposits, ideal-gas back-reaction, a relaxed particle timestep); 4c adds dust self-gravity through the existing Poisson solve; 4d builds the science configuration; Phase 5 takes particles onto refined meshes.

1.2 What Phase 4a must establish

1. **Dust fidelity.** The dust fork ships its validation in `../athenak_dust_tests/` (notebooks with printed reference values). On the merged tree, those references must reproduce: the damping ladder (second order in Δt , momentum to round-off), the NSH drift equilibrium (round-off over ten orbits), the damped dusty wave (spatial convergence), the random-placement and single-block layout checks, and the regression of the ion-neutral C-shock that shares the `imex2+` coefficient code.
2. **Gravity regression.** Three static multigrid tests whose numbers are documented in the multigrid validation document: the slab-open sheet test, the Tomida & Stone §4.1 contraction, and the boundary-annuli shearing wave including its zero-phase bit-identity with the periodic twin.
3. **Rank independence.** Dust and gravity at 1, 2 and 4 MPI ranks.
4. **Integrator.** The dust module’s IMEX coupling requires `imex2+`; all gravity validation so far used `rk2`. A gas-only stratified gravito-turbulence box under both integrators, compared column by column.
5. **The second merge (§4).** After catching up with the dust fork: the battery again, and the new PC2/hybrid couplings under `rk2`.
6. **Phase 4b (i) (§5).** The shear-periodic deposit fold: exactness, convergence, conservation, rank independence, and its effect on a structured 3D run.
7. **Phase 4b (ii) (§6).** The ideal-gas energy update: momentum dynamics unchanged, energy budgets of both modes versus the step, the heating rate against the analytic dissipation.

2 Test setup

2.1 Builds

Serial and MPI builds in the tree’s `build/` and `build-mpi/`, both with `-D PROBLEM=built_in_pgens -D Athena_ENABLE_FFT=ON`, the MPI one with `-D Athena_ENABLE_MPI=ON` (Open MPI via Homebrew). Everything below was first run on a scratch clone at the same commits and then rerun with the tree’s own binaries; the two sets of results are identical, including the gravito-turbulence histories bit for bit.

2.2 Inputs and commands

Dust inputs come from the dust-tests repository (paths relative to the directory holding both trees):

```
# damping ladder (tlim = 1, so dt = 1/ncycle)
for c in 0.4 0.2 0.1 0.05; do
  athena -i athenak_dust_tests/inputs/dust_damping.athinput particles/ppc=3 time/
    cfl_number=$c
done
# layout variants
athena -i .../dust_damping.athinput particles/ppc=4 problem/lattice=false
athena -i .../dust_damping.athinput particles/ppc=4 problem/lattice=false dust/deposit=
  ngp
athena -i .../dust_damping.athinput particles/ppc=3 meshblock/nx1=32 meshblock/nx2=32
# NSH equilibrium, dusty wave, C-shock regression
athena -i .../dust_nsh.athinput
athena -i .../dusty_wave.athinput mesh/nx1=$n meshblock/nx1=$((n/2))      # n = 32..256
athena -i inputs/ion-neutral/perp_cshock.athinput time/integrator=imex2+ time/nlim=500
# rank independence (drifting particles cross MeshBlock and rank boundaries)
mpirun -np $np athena -i .../dust_damping.athinput particles/ppc=3 problem/v0=0.5 dust/
  taus_3=0.5
# gravity statics
athena -i inputs/tests/mg_poisson_slab.athinput
athena -i inputs/tests/ts41_sin3.athinput
athena -i inputs/tests/mg_poisson_shear_bndry.athinput problem/profile=shwave problem/
  time0=0.37
# integrator twin (gas only, uniform grid)
athena -i inputs/shearing_box/gravito_turb_amr_test.athinput mesh_refinement/refinement=
  none \
  time/tlim=3.0 [time/integrator=imex2+ time/cfl_number=0.3]
```

Each dust problem generator writes a <basename>-errs.dat line per run; the gravity statics print their metrics to standard output. All result files are archived under validation/run/dust4a/ (see its README.txt), which the figure scripts figures/plot_dust4a_*.jl read.

3 Results

3.1 Dust fidelity

Damping ladder (Krapp et al. 2024 §3.1). Uniform gas at rest coupled to three dust species, $t_s = (0.01, 0.1, 1.0)$, $\epsilon_s = 1/3$ each, one particle of every species at every cell center so the discrete update per cell is the integrator's map of the ODE system. Errors are against a fine-step RK4 reference; the first run has $\Delta t > t_{s,1}$ (stiff regime).

	Δt	err _u	err _{v1}	err _{v2}	err _{v3}	$ \Delta P_{\text{tot}} $
	1.235×10^{-2}	1.3529×10^{-7}	1.3711×10^{-7}	1.4641×10^{-7}	6.8939×10^{-7}	1.6×10^{-16}
	6.211×10^{-3}	3.4627×10^{-8}	3.5096×10^{-8}	3.7325×10^{-8}	1.7630×10^{-7}	7.5×10^{-17}
	3.115×10^{-3}	8.7724×10^{-9}	8.8920×10^{-9}	9.4213×10^{-9}	4.4631×10^{-8}	2.0×10^{-16}
	1.558×10^{-3}	2.2075×10^{-9}	2.2377×10^{-9}	2.3649×10^{-9}	1.1225×10^{-8}	1.7×10^{-16}
order of err _u		1.98, 1.99, 1.99				
reference (dust tests, 2026-07-14)		err _u = 1.3529×10^{-7} , 3.4627×10^{-8} , 8.7724×10^{-9} , 2.2075×10^{-9} ; $ \Delta P \sim 10^{-16}$				

Every reference digit reproduces (Fig. 1). The layout variants: random positions with TSC and with NGP deposit give $\text{err}_u = 5.29 \times 10^{-6}$ and 5.77×10^{-6} (the 10^{-5} -level departure from the homogeneous ODE is physical: density fluctuations couple to pressure) with momentum to 3.8×10^{-16} ; the single-MeshBlock lattice run reproduces the four-block ladder row to all digits, which is the halo-exchange exactness check.




Figure 1: The damping ladder on the merged tree: second-order convergence of all four mean-velocity errors in Δt under the IMEX coupling (filled), the gas error under the PC2 coupling with `rk2` after the second merge (open, §4), and the total gas+dust momentum change of every IMEX run against the 10^{-13} acceptance line (right).

NSH drift equilibrium (2D r - z , ten orbits). Gas plus three species ($t_s = 0.01, 0.1, 1.0$; $\epsilon = 0.2, 0.5, 0.3$) started at the exact multi-species steady drift solution, with the gas forced by the pressure-gradient mimic $2\Omega\eta v_K$. After 4495 cycles every deviation is at round-off:

	$ \delta u_x $	$ \delta u_\phi $	$ \delta v_{x,1} $	$ \delta v_{\phi,1} $	$ \delta v_{x,2} $	$ \delta v_{\phi,2} $	$ \delta v_{x,3} $	$ \delta v_{\phi,3} $
serial	1.3×10^{-16}	6.2×10^{-17}	6.9×10^{-17}	5.9×10^{-17}	1.2×10^{-16}	2.4×10^{-17}	9.4×10^{-17}	4.3×10^{-17}
4 ranks	1.1×10^{-16}	9.0×10^{-17}	1.2×10^{-16}	2.8×10^{-17}	1.3×10^{-16}	1.0×10^{-17}	9.0×10^{-17}	3.3×10^{-17}

This is the test that certifies the rotation/forcing kicks and the implicit drag are composed unsplit inside the IMEX stages: any splitting error would show as secular drift.

Damped dusty sound wave (Laibe & Price 2011; Krapp et al. §3.2). The exact complex eigenmode of the linearized gas-dust system, L_1 errors of density and velocity against the analytic damped wave at $t = 2$:

N_x	cycles	$L_1(\delta\rho)$	$L_1(\delta v_x)$	order
32	320	1.836×10^{-8}	1.351×10^{-8}	—
64	321	6.896×10^{-9}	5.124×10^{-9}	1.41
128	641	1.837×10^{-9}	1.377×10^{-9}	1.91
256	1281	4.765×10^{-10}	3.572×10^{-10}	1.95

The code’s own root of the dispersion relation, $\omega = 4.529763 - 0.492158i$, is printed identically in every run. The order approaches two from below; the $32 \rightarrow 64$ step is polluted by the fixed-`cf1` time step at the coarsest resolution (the same cycle count at 32 and 64 cells), as the dust-tests notebook also observes.

3.2 Regressions

The ion-neutral C-shock (`inputs/ion-neutral/perp_cshock.athinput`, `integrator = imex2+`, 500 cycles) is the direct consumer of the refactored `imex2+` coefficient code: $\text{RMS-}L_1 = 2.071216 \times 10^{-3}$, bit-identical to the dust fork’s baseline. The Sod tube runs through the untouched hydro task list.

3.3 Rank independence, and the deadlock

The damping problem with drifting particles ($v_0 = 0.5$, $t_{s,3} = 0.5$, three lattice species) sends particles across MeshBlock and rank boundaries many times per run and exercises per-stage migration, the additive deposit exchange, the u^* ghost fill and the back-reaction exchange.

The first attempt at 4 ranks **hung at cycle 0**. The cause was in the merge itself: the dust fork’s additive deposit exchange posted per-neighbor MPI messages against the base class of July’s upstream, while this tree carries upstream #775, whose base-class receive posting is rank-packed (one aggregated message per peer rank). The sends had no receiver. After moving the exchange onto the rank-packed path (implementation document §4.4) with the serial path untouched (the damping ladder is bit-identical before and after the edit):

ranks	err_u	err_{v1}	err_{v2}	err_{v3}	$ \Delta P_{\text{tot}} $
1	1.158578×10^{-5}	1.187422×10^{-5}	1.490167×10^{-5}	6.153323×10^{-5}	5.2×10^{-15}
2	1.158578×10^{-5}	1.187422×10^{-5}	1.490167×10^{-5}	6.153323×10^{-5}	3.7×10^{-15}
4	1.158578×10^{-5}	1.187422×10^{-5}	1.490167×10^{-5}	6.153323×10^{-5}	3.6×10^{-15}
reference	1.158578×10^{-5}	(dust tests: “1-rank and 4-rank identical”)			

The tiny spread in $|\Delta P|$ is the reduction order of the diagnostic itself. The NSH run at 4 ranks is in the table of §3.1.

3.4 Gravity regression

test	metric	measured	documented
<code>mg_poisson_slab</code> , sheet profile (absolute, no mean subtraction)	max rel / L_2 rel	1.203474×10^{-12} / 3.078736×10^{-13}	1.2×10^{-12}
<code>ts41_sin3</code> (Tomida & Stone §4.1)	defect ratio per V-cycle	0.1254, 0.1249, 0.1249, 0.1249	0.125
<code>mg_poisson_shear_bndry</code> , shwave, $t_0 = 0.37$	max rel / L_2 rel	5.950118×10^{-3} / 2.057024×10^{-3}	5.95×10^{-3} / 2.057×10^{-3}
same, $t_0 = 0$, shear-periodic vs periodic- x_1 twin	both runs	5.506611×10^{-3} / 1.281655×10^{-3} , identical	bit-identical

At 4 ranks the slab and boundary-annuli solves reproduce the serial digits exactly, as documented for Phases 2b and 3.

The dust module is inert without a `<dust>` block by construction (the hydro and gravity paths do not see it), and these numbers confirm the shared infrastructure it touched — driver coefficients, boundary-values base class, output tables, physics registration — left the solver untouched.

3.5 The integrator twin

The Phase-3 small stratified gravito-turbulence box (`inputs/shearing_box/gravito_turb_amr_test.athinput:16H×16H×12H`, $64×64×48$, $\beta = 10$, slab-open multigrid gravity, refinement switched off) evolved to $t = 3\Omega^{-1}$ with `rk2` and with `imex2+`, both at `cfl_number = 0.3`. Relative differences of the history columns at $t = 3$:

	mass	E_{tot}	E_{int}	$\langle \rho c_s \rangle$	$\langle \rho w_{\text{grv}} \rangle$	$\langle \rho w_{\text{rey}} \rangle$	$\langle \rho \delta v^2 \rangle$
<code> imex2+ - rk2 / rk2 </code>	2.0×10^{-6}	7.0×10^{-5}	6.3×10^{-5}	3.5×10^{-5}	5.3×10^{-3}	3.2×10^{-4}	3.2×10^{-4}

Two different second-order integrators (three explicit stages of which the first advances by $\gamma\Delta t \approx 1.71\Delta t$, versus two) on a chaotic flow: agreement at 10^{-4} – 10^{-3} in the bulk quantities, with the small gravitational stress — a difference of large terms at this early time — the loosest at 5×10^{-3} (Fig. 2). The spikes of the relative difference of the stress columns at $t \approx 1.75$ and 2.5 in the figure are zero crossings of the stresses themselves (top-left panel), where a relative measure carries no information. This is the level at which `rk2`-validated gravity results transfer to dust runs; the full-size SC14 statistics under `imex2+` are a cluster item and will be attached to the next β -scan submission.

4 The second merge: regression, and the PC2/hybrid couplings

The dust fork’s `dust` branch had moved on by seven commits (to `1f3e1ed9`) while the first merge was based on `d6de1296`; commit `78320aee` merges them (implementation document §4.5). They bring the fork’s own rank-packed deposit exchange (taken in place of ours), CPU particle-mesh optimizations, and two new coupling modes selectable with `<dust>/coupling`: `pc2`, a per-cycle predictor–corrector particle update under `rk2`, and `hybrid`, which switches between PC2 and a first-order split backward-Euler branch on stiffness measures $\zeta = \max \Delta t/t_s$ and $\chi = \max \Delta t \rho_d/(\rho_g t_s)$.

4.1 Regression

Every table of §3 was rerun on the rebuilt serial and MPI binaries. The damping ladder, its three layout variants, the dusty-wave ladder, the C-shock, the drifting-particle test at 1, 2 and 4 ranks, and all multigrid statics (serial and 4 ranks) reproduce the printed digits above exactly; the NSH deviations move within their 10^{-16} round-off (the new non-atomic deposit changes the summation order); and the gas-only gravito-turbulence run is **bit-identical** to the pre-merge history, as it must be — the merge touched nothing on the gas path.

4.2 PC2 coupling under rk2

The inputs are the same test inputs with `integrator = rk2` and `coupling = pc2` added to the `<dust>` block (archived under `run/dust4a/merge2/`; the key has to be in the file because the command line cannot add parameters). Damping ladder:




Figure 2: The gas-only gravito-turbulence box under `rk2` (solid) and `imex2+` (dashed): stresses normalized by $\langle \rho P \rangle$, rms velocity fluctuation, internal and total energy, and the relative difference of the history columns.

Δt	err _u	err _{v1}	err _{v2}	err _{v3}	$ \Delta P_{\text{tot}} $
1.235×10^{-2}	2.4349×10^{-7}	2.2596×10^{-7}	2.4748×10^{-7}	1.2039×10^{-6}	2.4×10^{-16}
6.211×10^{-3}	5.2577×10^{-8}	4.8577×10^{-8}	5.3355×10^{-8}	2.5966×10^{-7}	4.5×10^{-17}
3.115×10^{-3}	1.2351×10^{-8}	1.1216×10^{-8}	1.2461×10^{-8}	6.0729×10^{-8}	5.0×10^{-17}
1.558×10^{-3}	2.9941×10^{-9}	2.7081×10^{-9}	3.0167×10^{-9}	1.4707×10^{-8}	5.1×10^{-17}
order of err _u	2.23, 2.10, 2.04	(IMEX ladder: 1.98, 1.99, 1.99)			

Second order, converging onto the Δt^2 slope from above, with errors about 1.4–1.8× the IMEX ones at equal Δt and momentum at round-off (Fig. 1, open markers). The dusty wave:

N_x	$L_1(\delta\rho)$	$L_1(\delta v_x)$	order	(IMEX order)
32	1.612×10^{-8}	1.182×10^{-8}	—	—
64	4.483×10^{-9}	3.310×10^{-9}	1.85	1.41
128	1.139×10^{-9}	8.404×10^{-10}	1.98	1.91
256	2.858×10^{-10}	2.109×10^{-10}	2.00	1.95

Cleaner second order than the IMEX ladder and smaller errors at every resolution. The NSH equilibrium under pc2 holds to round-off over the ten orbits (3.2×10^{-16} , 1.3×10^{-16} , 3.1×10^{-16} , 2.1×10^{-16} , 3.6×10^{-16} , 2.4×10^{-16} , 1.0×10^{-17} , 1.7×10^{-16} for the eight components of §3.1). Rank independence with the drifting particles:

ranks	err _u	err _{v1}	err _{v2}	err _{v3}	$ \Delta P_{\text{tot}} $
1	6.805779×10^{-6}	5.812361×10^{-6}	7.380154×10^{-6}	3.360985×10^{-5}	1.5×10^{-14}
2	6.805779×10^{-6}	5.812361×10^{-6}	7.380154×10^{-6}	3.360985×10^{-5}	1.9×10^{-15}
4	6.805779×10^{-6}	5.812361×10^{-6}	7.380154×10^{-6}	3.360985×10^{-5}	3.9×10^{-15}

4.3 Hybrid coupling and the split-BE branch

Two runs characterize the fallback. The three-species NSH test under `coupling = hybrid` (automatic selection) reports `pc2_cycles=0 split_be_cycles=4494` with $\zeta = 0.549$: the $t_s = 0.01$ species puts $\Delta t/t_s$ just above the 0.5 entry threshold, so the run stays in the first-order branch throughout, and the equilibrium is no longer an exact fixed point — deviations of $1\text{--}12 \times 10^{-5}$ after ten orbits instead of round-off. This is the designed behavior, not a defect: the first-order splitting error of a stiff fallback, on a problem whose IMEX and PC2 answers are exact. The stiff damping case of the dust tests ($t_s = 10^{-4}, 10^{-3}, 2 \times 10^{-3}$, $\Delta t/t_s$ up to 125) under `hybrid` runs split-BE for all 81 cycles ($\zeta = 32.8$, $\chi = 12.6$) and lands on the exact relaxed state (errors $\leq 4.9 \times 10^{-16}$, momentum 1.1×10^{-16}), the L-stability that the branch exists for. For production runs in the resolved regime the recommendation is therefore `coupling = pc2`; `hybrid` is the safety net for configurations with stopping times below the timestep.

5 Phase 4b (i): the shear-periodic deposit fold

The mechanism is in the implementation document §5: the x_1 ghost slabs of the shear-face Mesh-Blocks are gathered into global y – z planes, remapped in y by $\mp y_{\text{sh}}$ (the adjoint of the ghost-fill signs), and added into the active edge strips across the face, while the unsheared plain-periodic x_1 buffers of those blocks are skipped. Data: `run/dust4b/` (README there).

5.1 The unit test

`pgen_name = dust_deposit_shear` (`inputs/dust/dust_deposit_shear.athinput`: $32 \times 32 \times 8$, $16 \times 16 \times 8$ blocks so both x_1 columns are face blocks, `nghost` = 3, $q = 3/2$, $\Omega = 1$, $L_x = 1$). Three fields live only in the x_1 ghost slabs of the face blocks: a smooth periodic g on the slab rows within the block's own y/z range, a constant on the same rows, and a constant on the whole slab including the corner ghosts (whose expectation is the analytic overlap multiplicity).

case	$y_{\text{sh}}/\Delta y$	max rel (g)	const	multipl.	leak	mass
plain pass only (<code>fold = false</code>)	0	0	0	0	0	(none delivered)
fold, <code>time0 = 0</code>	0	0	0	0	0	5.9×10^{-16}
fold, 3-cell integer shift	3	3.5×10^{-16}	0	0	0	5.9×10^{-16}
fold, generic phase, $4 = 2 = 1$ ranks	35.52	1.332091×10^{-3}	0	(frac.)	0	10^{-15}

The first row is the suppression check: with the fold switched off the receiving strips stay *exactly* zero, so no unsheared x_1 contribution survives (the first version of the direction table failed this row for the $x_3 x_1$ corners; implementation document §5.4). Rows two and three are the exactness checks, including the multiplicity of overlapping corner rows. At a fractional shift the multiplicity field is a step profile and carries an $O(1)$ remap error at block boundaries, as expected; the smooth field's error is the remap interpolation error, which the last row shows to be identical at 1, 2 and 4 ranks to every printed digit.

Convergence. At a fixed fractional shift ($y_{\text{sh}}/\Delta y = n_y/2 + 1/2$, so `eps` = 1/2 at every resolution; the generic-phase scan is confounded by the changing fraction):

n_y	DC L_1	DC max	PLM L_1	PLM max	PPMX L_1	PPMX max
32	1.86×10^{-3}	1.24×10^{-2}	2.69×10^{-4}	6.40×10^{-3}	1.11×10^{-4}	2.11×10^{-3}
64	4.66×10^{-4}	3.18×10^{-3}	3.89×10^{-5}	1.33×10^{-3}	9.16×10^{-6}	4.76×10^{-4}
128	1.17×10^{-4}	7.96×10^{-4}	5.33×10^{-6}	6.49×10^{-4}	9.90×10^{-7}	1.17×10^{-4}
256	2.91×10^{-5}	1.99×10^{-4}	6.21×10^{-7}	1.47×10^{-4}	9.78×10^{-8}	2.92×10^{-5}
order	2.00	2.0	2.8–3.1	~ 2	3.2–3.6	2.0

Donor cell is exactly second order (the half-cell shift is symmetric); PLM, the default (matching the u^* remap), is second order in the maximum norm and better than second in L_1 ; PPMX is third order in L_1 with a limiter-bound maximum (Fig. 3). Mass is conserved to $\leq 3 \times 10^{-15}$ and the leak is exactly zero in every run.

5.2 Regression: the 3D shear-periodic NSH equilibrium

The new tracked input `inputs/dust/dust_nsh_shear3d.athinput` ($32 \times 32 \times 8$, $2 \times 2 \times 1$ blocks, ten orbits) is the azimuthally uniform state for which the old unsheared exchange was already exact:

	max deviation, before the fold	with the fold	PC2, with the fold	4 ranks
NSH, 3D shear-periodic	2.7×10^{-16}	2.6×10^{-16}	1.8×10^{-15}	1.1×10^{-16}




Figure 3: Convergence of the shear-periodic deposit fold at a fixed half-cell fractional shift, for the three remap orders.

5.3 A structured run: rank independence and the size of the correction

Random particle placement in the NSH generator is rank-dependent by construction (it hashes tags to *local* block indices), which made a first serial-vs-MPI comparison look like a defect; with test particles the same runs agreed to 10^{-13} , and the cycle-0 dust density already differed. The deterministic alternative is the lattice with a 50% sinusoidal azimuthal modulation of the particle masses (`problem/mass_mod_amp = 0.5`, `check_equilibrium = false`), IMEX coupling with back-reaction, one orbit. Relative differences of the gas x -kinetic energy and of the species-1 radial-velocity sum:

t	serial vs 2 ranks	serial vs 4 ranks	fold vs. no fold	PLM vs. PPMX remap
0.81	$10^{-14} / 7 \times 10^{-15}$	$5 \times 10^{-15} / 9 \times 10^{-15}$	$4.2 \times 10^{-2} / 2.7 \times 10^{-2}$	$3.7 \times 10^{-7} / 7.9 \times 10^{-8}$
3.20	$1.2 \times 10^{-14} / 5 \times 10^{-15}$	$1.2 \times 10^{-14} / 9 \times 10^{-15}$	$9.7 \times 10^{-3} / 5.8 \times 10^{-3}$	$5.4 \times 10^{-9} / 2.1 \times 10^{-8}$
4.80	$6 \times 10^{-15} / 7 \times 10^{-15}$	$6 \times 10^{-15} / 5 \times 10^{-15}$	$6.8 \times 10^{-2} / 3.9 \times 10^{-2}$	$4.6 \times 10^{-7} / 3.1 \times 10^{-7}$
6.28	$6 \times 10^{-15} / 5 \times 10^{-15}$	$6 \times 10^{-15} / 5 \times 10^{-15}$	$5.4 \times 10^{-3} / 3.6 \times 10^{-3}$	$4.4 \times 10^{-7} / 2.6 \times 10^{-7}$

Three conclusions. The fold is rank-independent to round-off on a structured field with back-reaction. The old unsheared exchange (`dust/shear_fold = false`) was a 1–7% error in the gas kinetic energy for a structured state, oscillating with the shear phase — not a small correction. And the remap order is a 10^{-7} effect on this run, so PLM is converged for the purpose.

6 Phase 4b (ii): the ideal-gas energy update

Mechanism: implementation document §6. Every drag kick on the gas momentum now books the matching kinetic-energy change (internal energy fixed under the kick); with `dust/drag_heating = true` the frictional dissipation of each particle kick is deposited as heat. Data: `run/dust4b/energy/`.

6.1 The damping test with an ideal gas

`inputs/dust/dust_damping_ideal.athinput`: the three-species damping problem of §3.1 with `eos = ideal`, $\gamma = 1.4$ and $p_{\text{gas}} = 10^{-5}$ (comparable to the 1.4×10^{-5} of kinetic energy the drag dissipates, so the six-digit history resolves the budget; the first attempt with $p = 1$ could not), `dtmax` = 1/95 (the isothermal ladder’s first step). The internal energy is $e = E - K_{\text{gas}}$ from the hydro history, the dust kinetic energy $K_{\text{dust}} = \sum_s \frac{1}{2} \epsilon_s \rho_0 V \bar{v}_s^2$ from the species velocity history (exact on the lattice).

Momentum dynamics. The four errors of the ODE test are identical between the default and heating modes to all printed digits (9.733252×10^{-8} etc. under IMEX, 1.660169×10^{-7} under PC2) — pressure and heat do not touch the velocities on a uniform state — and momentum is conserved to 10^{-16} .

Energy budgets versus the step. Errors relative to the dissipated kinetic energy, default mode ($|\Delta e|$) and heating mode ($|\Delta(E_{\text{gas}} + K_{\text{dust}})|$), with the order between successive rows in parentheses:

Δt	IMEX default	IMEX heating	PC2 default	PC2 heating
1/95	1.07×10^{-2}	3.81×10^{-2}	2.15×10^{-2}	5.89×10^{-2}
1/190	6.83×10^{-3} (0.65)	2.80×10^{-2} (0.44)	7.14×10^{-3} (1.59)	3.09×10^{-2} (0.93)
1/380	4.17×10^{-3} (0.71)	1.93×10^{-2} (0.54)	2.82×10^{-3} (1.34)	1.53×10^{-2} (1.01)
1/760	2.43×10^{-3} (0.78)	1.23×10^{-2} (0.65)	1.24×10^{-3} (1.18)	7.56×10^{-3} (1.02)

The default-mode residual is the register-combination gap $\gamma_0 \gamma_1 |\mathbf{K}_1|^2 / 2\rho$ per step (implementation document §6.3): the 1% at $\Delta t = 1/95$ is reproduced by that estimate with the stage-1 kick of the $t_s = 0.01$ species, and the sub-first-order IMEX slope is the saturation of that kick while $\Delta t \gtrsim t_s$ (the orders climb toward one as the step shrinks; PC2, with its RK2 weights, is already first order and better). The heating mode inherits the same gap on the dust side and converges at first order. Both are truncation errors of the tableau, not bookkeeping mistakes: the sign, size and Δt -scaling are the predicted ones.

6.2 The NSH state with an ideal gas

`inputs/dust/dust_nsh_shear3d_ideal.athinput` ($p = 1$, two orbits):

	max deviation of the drift equilibrium	$\dot{e}$ / analytic dissipation rate
default mode, <code>cf1 0.45</code>	4.6×10^{-15}	(no heat: e constant to 8×10^{-5})
heating, <code>cf1 0.45</code> ($\Delta t/t_{s,1} = 1.16$)	4.6×10^{-15}	1.0148
heating, <code>cf1 0.225</code>	—	1.0074
heating, <code>cf1 0.1125</code>	—	1.0037
heating, <code>cf1 0.45</code> , 4 ranks	6.7×10^{-16}	1.014777 (serial: 1.014777)

The analytic rate is $\sum_s \epsilon_s \rho_0 |\mathbf{u} - \mathbf{v}_s|^2 / t_s$ from the equilibrium velocities. The heating rate converges to it at first order (the error halves with the step; the stiff $t_s = 0.01$ species dominates), the equilibrium is untouched by either mode, and 4 ranks reproduce the serial rate to every digit.

6.3 Regression

Isothermal runs cannot enter the new branches. The 2D damping ladder is bit-identical to the second-merge battery; the 2D and 3D NSH errors and the structured-run histories differ from their Phase-4b (i) values only in quantities at the 10^{-16} level (the deposit kernels changed shape and the compiler schedules the fused multiply-adds differently), and the 4-rank drifting-particle test reproduces the reference 1.158578×10^{-5} with the now four-channel exchange.

7 Summary

Phase 4a delivers the dust module inside the multigrid tree with both sides certified unchanged: the dust fork’s acceptance numbers and the gravity solver’s documented statics reproduce to every printed digit, serially and at 4 ranks, and the one merge defect (the additive deposit exchange against rank-packed messaging) is fixed with a rank-independence table to show for it. The integrator the dust module’s IMEX coupling imposes, `imex2+`, tracks `rk2` on the stratified gravito-turbulent box at the level two second-order schemes should — and after the second merge that constraint is optional: the PC2 coupling runs under `rk2` at second order with the same conservation properties, and is the recommended production path in the resolved drag regime.

Phase 4b (i) closes the first gap Phase 4a left open: ghost deposits now cross the shear-periodic radial faces with a conservative azimuthal remap, certified exact at integer shifts, convergent otherwise, conservative, leak-free and rank-independent, and shown to matter at the percent level on structured states. Phase 4b (ii) makes back-reaction usable with an ideal gas: the energy book-keeping is exact per kick, the residuals are the RK combination gap (a percent of the dissipated energy at a step equal to the stiffest stopping time, converging), and the optional frictional heating reproduces the analytic dissipation rate to first order. Still open in 4b: a particle timestep that does not pay the full-shear-velocity penalty in wide boxes, tested on the NSH state under the relaxed step.